



**University of
Nottingham**

UK | CHINA | MALAYSIA

Tabletop War Game Assistant

Submitted April 2024, in partial fulfilment of
the conditions for the award of the degree:
BSc Hons Computer Science

20363169
School of Computer Science
University of Nottingham

I hereby declare that this dissertation is all my own work, except as indicated
in the text:

Signature: NB
Date: 13/04/24

I hereby declare that I have all necessary rights and consents to publicly
distribute this dissertation via the University of Nottingham's e-dissertation
archive.*

Table of Contents

1. Abstract	1
2. Introduction and Project Intention	1
3. Background	2
3.1. game board Components	2
3.2. Gameplay	4
3.2.1. Shooting and Line of Sight	4
3.2.2. Key Takeaways	6
3.3. Community	6
4. Related Work	7
4.1. Previous Mixed Reality Tabletop Systems	7
4.2. Object Detection and Tracking Technologies	10
4.2.1. RFID	10
4.2.2. Machine / Deep Learning	11
4.2.3. ARKit	11
4.2.4. Computer Vision	12
4.3. Fiducial Markers and Tag Detection	12
5. Project Description	13
6. Methodology and Design	14
6.1. Tracking Methodology	14
6.1.1. Camera Setup	14
6.1.2. Terrain Detection	15
6.1.3. Operative Detection and Identification	16
6.2. Game Board Representation	18
6.2.1. Terrain and Operative Placement	18
6.2.2. GUI design	19
6.2.3. Line of Sight	19
7. Implementation	20
7.1. Model Tracking	20
7.1.1. Design	20
7.1.2. Video feed	21
7.1.3. Calibration	22
7.1.4. Homography	23
7.1.5. Board Detection	25
7.1.6. Tag Detection	26
7.1.7. Model Identification	27
7.1.8. Optimisations	28
7.2. Terrain Detection	29
7.3. Game Board Representation	30
7.3.1. Linking of detected models and terrain to the virtual board	30
7.3.2. Line of Sight	31
8. Evaluation	33
8.1. Tracking System	33
8.1.1. Accuracy	34
8.1.2. Multiple Tag Tracking	37
8.1.3. Tracking With Terrain	37
8.1.4. Summary	40
9. Summary and Reflections	40

9.1. Project Management	40
9.2. Future Work and Reflections	41
9.2.1. Order Tokens	41
9.2.2. Odds to Hit	42
9.2.3. Increasing total number of operatives	42
9.2.4. Server Client Architecture For Game Board and Detection	42
9.2.5. Solo Play	42
9.2.6. Detection Upgrades	42
9.3. LSEPI	43
9.3.1. Accessibility	43
9.3.2. Open Source	43
10. Final Reflections	43
11. Bibliography	44

1. Abstract

This project aims to create a computer vision based tracking system for physical tabletop miniatures for an example wargaming system. As wargaming rules are quite complex, the system will also aim to create a viable representation of certain rules to aid the players. This involves the creation of custom tags to track miniatures aimed to handle occlusion, provide accurate information whilst also not requiring external hardware.

2. Introduction and Project Intention

Tabletop war-gaming is a popular hobby which has recently seen a surge in popularity during the COVID-19 pandemic, however, with a high barrier to entry, it remains niche and inaccessible to many. The rules to tabletop war-games can be complex and difficult to learn. This can be daunting for new players, putting them off the hobby as well as causing disagreement between seasoned players over rules interpretations.

Some of the most popular war-gaming systems are produced by Games Workshop [1]. One of their more popular systems, Warhammer 40k, has a core rule-book of 60 pages [2] and the simplified version of another game system, Kill Team, is a rather dense three page spread [3]. This complexity can be off putting to new players. Many tabletop / board games suffer from a player's first few games taking significantly longer than is advertised due to needing to constantly check the rules.

As well as this, new players are likely to take longer to make decisions as they are not familiar with the game's mechanics of what constitutes a "good" move¹. This can be exacerbated by the game's reliance on dice rolls to determine the outcome of actions. Meaning that seemingly "optimal" moves do not always result in favourable outcomes, causing an extended learning period for an already complex game.

Kill Team is a miniature war-game known as a "skirmish" game. This means the game is played on a smaller scale with only ~20 miniatures on the table at one time. The aim of the game is to compete over objectives on the board for points. Each player takes turns activating a miniature and performing actions with it. The game uses dice to determine the results of your models engaging in combat with each other with the required rolls being determined by the statistics of each "operative" (a miniature) involved and the terrain.

¹sometimes referred to as 'analysis paralysis'



Figure 1: An example of the Kill Team tabletop game using the Gallowdark terrain. The terrain we will focus on here are the flat walls and pillars. [4]

Video games help on-board new players by having the rules enforced by the game itself. For example, a digital chess game will tell you when you are in check. A physical chess game requires players to recognise when they are in check themselves. This project aims to bring a similar methodology to tabletop war-gaming, specifically the Kill Team Lite [3] system using the Gallowdark setting. The Kill Team Lite rules are publicly available from Games Workshop's website and are designed to be played on a smaller scale to other war games, making it a good candidate for a proof of concept. As well as this, the Gallowdark [5] setting streamlines the terrain used and removes verticality from the game, making implementation much simpler.

Developing a system that can digitally represent a physical Kill Team board would allow for the creation of tools to assist players. For example, a digital game helper would remove the burden of rules enforcement from the players and onto the system, allowing players to focus on the game itself or allow players to make more informed game decisions by being able to preview the options available to them. This could also be utilised to record a game and view a replay or be used for content creation to make accurate board representations to viewers with digital effects.

This project will focus on the development of a system that can track the position of miniature models and terrain pieces on a Kill Team board which can subsequently be displayed digitally. From here, we aim to implement proof of concept visualisation of a few select game rules on the digital board to demonstrate that the tracking system provides the necessary information to process said rules.

3. Background

To determine the specific goals for this project, it is important to provide some more context on the gameplay and community surrounding Kill Team. As this is a very complex game, we will be omitting a large percentage of the rules, focusing on the topics that will have an impact on the overall design or provide unique challenges.

3.1. game board Components

Before looking at the game rules, we will first break down the physical components of the game board to determine what needs to be tracked and problems that may arise from this.

A game of Kill Team is composed of two players, each with a group of “operatives” (miniatures) referred to as a “Kill Team”. Each Kill Team has its own unique rules, with each operative having its own set of unique statistics and special abilities. A miniature is composed of a circular base with a model placed on top. The diameter of the base is either 25mm, 28.5mm, 32mm, 40mm or occasionally 50mm.



Figure 2: An example of a Kasrkin Kill Team operative on a 28.5mm base [6]. The height of the operative is ~35mm excluding the base. This project will focus on getting the system functional with the Kasrkin models on 28.5mm bases.

It is worth noting that it is common for models to extend past the edge of their base as seen in Figure 2. Whilst the example above is somewhat minor in its overlap, different models can be more extreme in their extension. As a result of this the system will need to be able to locate and identify an operative from only a partial view of the base, or its surroundings, from a bird’s eye view. The detection system will either need to be above the board if using visual detection methods or below the board if using, for example, an RFID method. This is due to the terrain potentially surrounding an operative, so having a camera on each side of the board will not guarantee it is visible as can be seen in Figure 1.

The game is played on a 30” by 22” board, referred to as a “Killzone”. The Gallowdark ruleset instead uses a 27” by 24” board with the specialised terrain shown in Figure 1. As stated previously, the terrain we are focusing on here are the thin walls with pillars connecting them. The terrain is all at the same height, but the operatives are shorter than the terrain. As a result, the system will need to find a solution to detect operatives behind terrain. From a birds eye view, the terrain will block part of the operative due to parallax. For this project, we will focus on using a half sized board to simplify this problem.

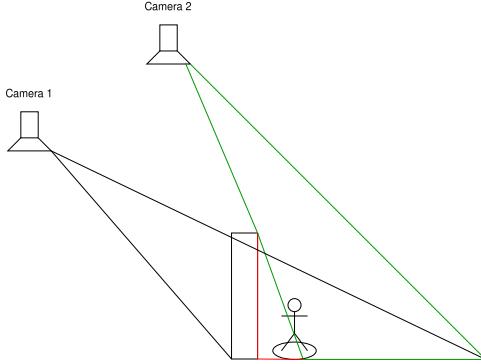


Figure 3: An example of the problem introduced by parallax. The camera is placed above the board offset from the operative. The operative is partially obstructed by terrain. The green triangle represents the parts visible to the camera. The red lines demonstrate the areas blocked by the terrain.

The further away the camera is from the operative on the x axis, the more the terrain will block the operative. However, as the camera travels away on the y axis, the terrain will block less of the operative. For a camera implementation this would also mean we lose quality in the image. As a result, a camera based system would need to find an ideal distance from the board which provides enough quality in the image, whilst also being able to see enough of an operative if it is obscured by terrain.

3.2. Gameplay

The gameplay of Kill Team focuses around a turn based system with a full game consisting of 4 rounds - called “turning points”. Each player takes turns “activating” a single operative and performing actions with it. The number of actions available to an operative is defined by its statistics². Once an operative has performed its maximum number of actions, play is handed to the other player. That same operative may not be “activated” until every other operative, still in play, on their team has been activated. Once every operative on both teams has been activated the next turning point begins resetting the activation of each operative still in play. Once all 4 turning points have been completed the game ends and a winner is decided.

This project is focussed on creating a digital representation of the board, and then implementing a few select rules from the game to demonstrate the system’s capabilities. The rules for “movement” and “shooting” present themselves as good candidates for this due to both their complexity and frequency within the game.

Please note that the explanations for these rules are abstracted from the Kill Team Lite rules published publicly by Games Workshop [3].

3.2.1. Shooting and Line of Sight

Operatives can be set to two states: “concealed” or “engaged”. These states are used to determine whether an operative is a valid target or whether it is able to make offensive actions, such as shooting. An operative’s state is denoted by a small triangle of orange card pointing to the model. The system will need a way to track the state of an operative, whether this be through detecting the orange markers or manually updating the state.

²There are a lot of exceptions to this with abilities being able to modify the stats of themselves or other operatives, but for the sake of simplicity we are going to ignore this and assume stats are set.

Anecdotally, the Kill Team line of sight rules tend to be the most complex part of the game and are constructed in such a way that experienced players can abuse them to gain an advantage such as one way shooting³.

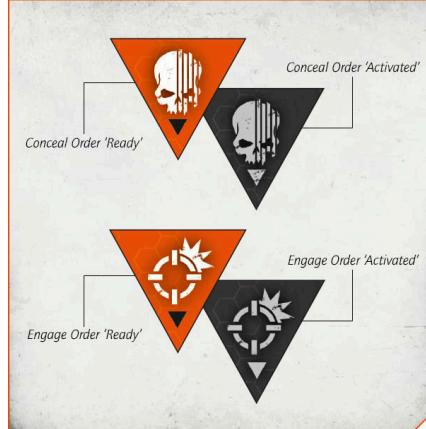


Figure 4: An example of the markers used to denote the state of an operative [7]. The board will need to be able to mark operatives as appropriate and display this to the user.

There are a set of requirements that must be met to determine whether an operative is able to be targeted by a shooting attack. We will use **attacker** to refer to the operative making the attack and **defender** to refer to the intended target.

For a defender in the engaged state:

1. The defender must be **visible**.
2. The defender must be not **obscured**.

For a defender in the concealed state:

1. The defender must be **visible**.
2. The defender must be not **obscured**.
3. The defender must be not **in cover**.

Kill Team defines **visible**, **obscured** and **in cover** very strictly. This is done using **cover / visibility lines** which are straight lines 1mm wide drawn from one point on an operative to another operative.

Visible is used to simulate whether a defender can be physically seen by an attacker

For a defender to be (visible) the following must be true:

1. A visibility line can be drawn from the head of the attacker to any part of the defender's **model**, not including the base.

Obscured is used to simulate whether a defender is blocked by large terrain, such as a wall, while also containing edge cases for operatives peeking round said terrain or attempting to use it directly as cover.

When drawing cover lines the attacker picks one point on their base and draws two cover lines from that point to every point on the defenders base, to create a cone. In practice this means drawing two cover lines to the extremes of the defender's base.

³This is a technique used to allow an attacker to be able to select a valid target but the defender is unable to fire back.

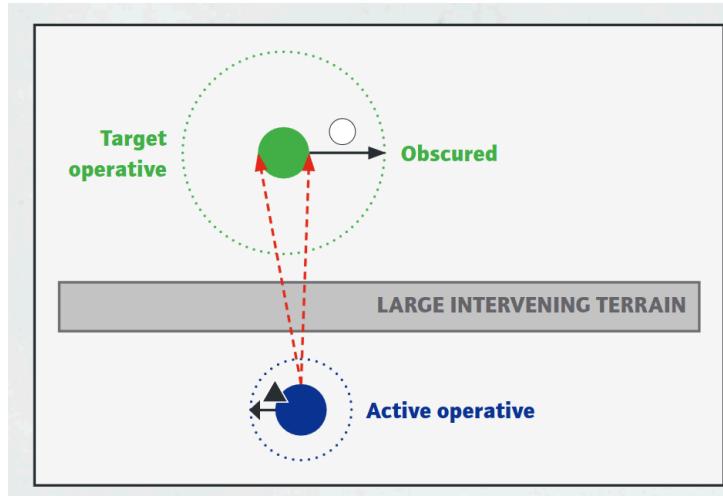


Figure 5: An example of cover lines [3].

For a defender to be **obscured** the following must be true:

1. The defender is $>2"$ from a point where a cover line crosses a terrain feature which is considered obscuring⁴.
1. If the attacker is $<1"$ from a point in which a cover line crosses a terrain feature, then that part of the feature is not obscuring.

In cover simulates whether an operative is attempting to intentionally take cover behind a piece of terrain. A concealed operative is hiding behind the terrain whereas an engaged operative is poking around / over it.

For a defender to be **in cover** the following must be true:

1. The defender is $>2"$ from the attacker
2. The defender is $<1"$ from a point where a cover line crosses a terrain feature.

3.2.2. Key Takeaways

The above rules outline what information the system will need to be able to track and display to the user. It also highlights the complexity of the rules to justify the reasonings behind this project.

1. It is important to know the positions of operatives and terrain.
2. The system will need to be able to know the size of an operative's base.
3. Rotation of operatives is not needed to be tracked (except for visibility when determining **visible**).
4. Operatives need to be marked as concealed or engaged.
5. Accuracy is important to the system. Visibility / cover lines being 1mm wide and distances measured in inches will require accurate measurements.
6. The system will need to abstract the above rules whilst also being able to display the reasoning behind the application. For example, if a defender is obscured the system should show what is obscuring and from what point the firing cone was drawn.

3.3. Community

Tabletop wargames require you to build and paint your own miniatures, because of this the target audience tends to be more of the creative type. Using this information we can allow some liberty in

⁴For our use case the terrain walls in Gallowdark are all considered obscuring terrain.

what they might need to do to utilise this system. For example, creating homemade tags is something that creatives might be more inclined to do as opposed to something more technical using RFID or infrared sensors. The whole system should be designed to be as accessible as possible to the target audience, requiring little to no specialised equipment.

The miniatures are also not defined in their appearance. They can be painted in any scheme, be customised to have completely different looks or use entirely different models than what is intended⁵. This means the tracking system needs to be ambivalent of what the model looks like, both in colour and silhouette.

Since miniatures are highly customisable players tend to be very attached to them. So one important requirement is to not be invasive to the miniature. This rules out requiring a certain paint job, placing stickers on heads or putting QR codes above operatives. The system should aim to obscure models as little as possible and cause no damage.

4. Related Work

4.1. Previous Mixed Reality Tabletop Systems

Some companies, such as The Last Game Board [8] and Teburu [9], sell specialist game boards to provide a mixed reality experience.

The Last Game Board achieves this through utilising the touch screen to recognise specific shapes on the bottom of miniatures to determine location and identity. The Last game board is 17" x 17", as a result the number of game systems which are compatible is limited. However, you can connect multiple systems together. The drawback of this is the price point for the system is rather high, with boards starting at ~\$350. It is also worth noting that this system has not received great reviews, with alpha users reporting: long load times, low FPS, graphical and sound glitches, lack of updates and even screens refusing to display half the screen for certain applications.



Figure 6: The Last Game Board touchscreen tabletop system [8]

Teburu [9] instead takes an RFID based approach, providing a base mat that allows you to connect squares containing RFID receivers and game pieces containing an RFID chip. Teburu connects to a tablet device to provide the digital experience as well as to multiple devices for individual player in-

⁵Known as proxying.

formation. Teburu games allow for game pieces to either be in predetermined positions or within a vague area i.e. within a room.

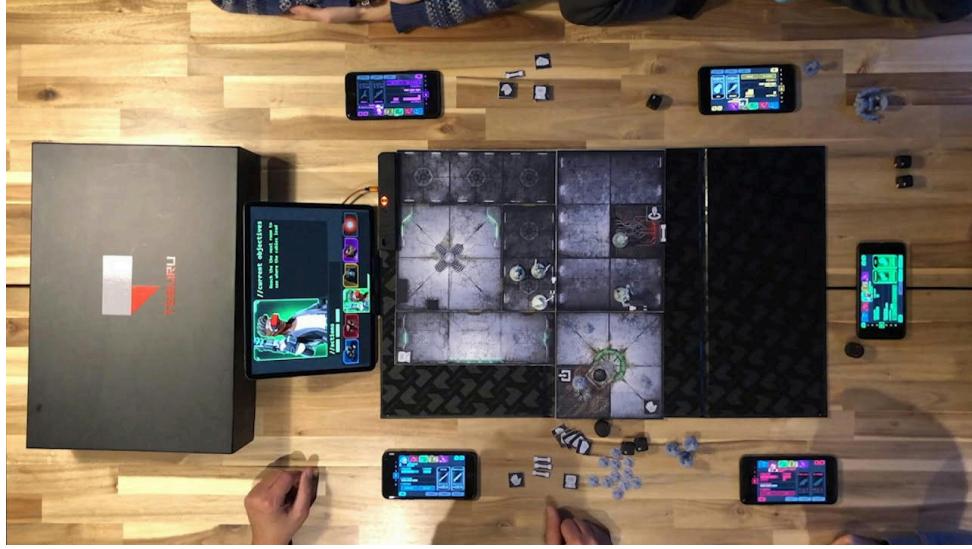


Figure 7: The Teburu Game System [10] showcasing The Bad Karmas board game. The black board is the main game board. The squares above connect to the board below to transmit the RFID reader information back to the system for display.

An RFID based approach is also used in “An RFID-based Infrastructure for Automatically Determining the Position and Orientation of Game Objects in Tabletop Games” (Steve Hinske and Marc Langheinrich [11]) which places an antenna grid below the game board to detect RFID chips within models. This allowed them to find what chip is in range of what antenna, allowing them to find the general location of a game piece. This worked particularly well for larger models where you could put RFID chips far away from each-other on the model. Using the known positions of the chips and dimensions of the model combined with which antenna said chips are in range of, allows you to determine an accurate position of each model. They also go into alternate RFID approaches which will be discussed later when outlining the chosen methodology.

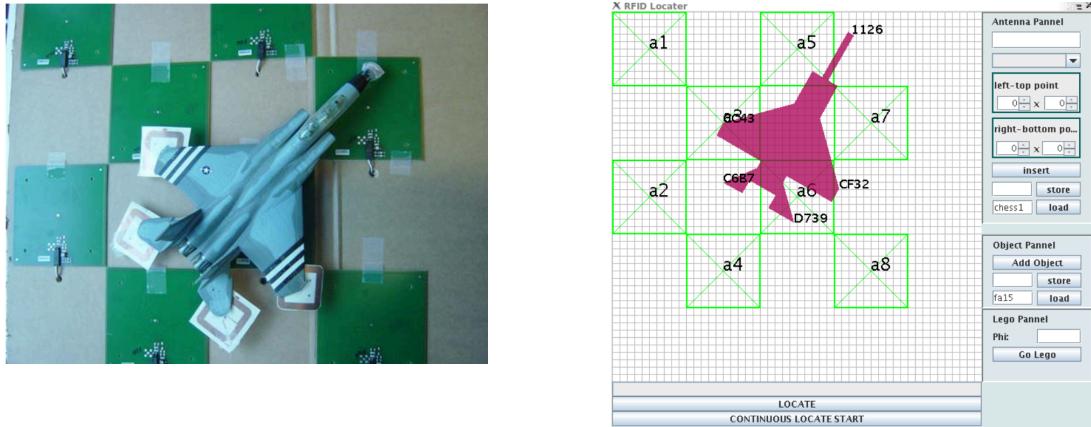


Figure 8: An example of Steve Hinske and Marc Langheinrich's approach depicting the antenna grid, RFID tags and physical model alongside the computer's prediction of the model's position

Surfacescapes [12] is a system developed in 2009 by a group of masters students at Carnegie Mellon as a university project. Surfacescapes uses a Microsoft Surface Tabletop (the product line was rebranded to PixelSense in 2011). This uses a rear projection display, and 5 near IR cameras behind the screen [13]. This allows the PixelSense to identify fingers, tags, and blobs touching the screen using the near

IR image. Surfacescapes utilises this tag sensing technology to track game pieces using identifiable tags on the bases of miniatures.



Figure 9: An example of surfacescapes' in use on the Microsoft Surface Tabletop [14]. The position of the models have been tracked by the system and outlined with a green circle.

Foundry Virtual Tabletop [15] is an application used to create fully digital Dungeons and Dragons tabletops. These can either be used for remote play or in person play using a horizontal TV as a game board. Foundry VTT allows for the creation of modules to add new functionality to your virtual tabletop. One such module is the Material Plane [16] module which allows the tracking of physical miniatures on a TV game board. This functions by placing each miniature on a base containing an IR LED with an IR sensor then placed above the board. This can be configured to either move the closest “virtual” model to where the IR LED is or (with some internal electronics in the bases) can be set up to flash the IR LED in a pattern to attach different bases to specific models. An indicator LED is present to show when the IR LED is active.



Figure 10: An example of one of the Material Plane bases. A miniature would be attached to the top. [17].

4.2. Object Detection and Tracking Technologies

Finding a method to detect and find the positions of miniatures on a game board, whilst not obstructing the game, is the main focus of this project. This section will outline the different technological approaches that could be taken to achieve this.

4.2.1. RFID

An RFID based approach appeared to be a good solution initially. This would involve embedding RFID chips underneath the bases of the miniatures. Then some method of reading these chips would then need to be embedded either within or underneath the game board. There are a number of different approaches that could be taken to locating RFID chips which have been outlined by Steve Hinske and Marc Langheinrich [11] in their work on a similar project.

RFID solutions would require either an antenna grid underneath the game board or multiple individual RFID readers. This would be a viable option as hiding an antenna grid below a board is a relatively simple and unobstructive task. The same goes for hiding RFID readers beneath a table.

The main drawback of RFID is that a reader (at its core) can only detect if a chip is in range or not (referred to as “absence and presence” results). Due to this, some extra methodology would need to be implemented to determine the position of a chip.

One approach utilises increasing the range of an RFID reader to its maximum and then subsequently reducing the range repeatedly. Using the known ranges of the readers it would be possible to deduce which tags are within range of which reader - and subsequently deduce from which ranges it is detectable in the position of the RFID chip. This approach could in theory provide a reasonably accurate position of the RFID chip. However, this approach would take a longtime to update as each RFID reader would need to perform several read cycles. Combined with problems caused from interference between the chips / readers, the fact that being able to vary the signal strength of an RFID reader is not a common feature and the need for multiple RFID readers, this approach does not meet the requirements for this project.

Another approach utilises measuring the signal strength received from an RFID chip and estimating the distance from the reader to the chip, known as received signal strength indication (RSSI) . This approach is much quicker than the previous method needing only a single read cycle to determine distance. Most modern day RFID readers can report RF phase upon tag reading. However, current RSSI methods have an error range of ~60cm caused by noise data, false positives / negatives and NLOS (non-line of sight) [18]. This won't work for this project given than the board size is 70 x 60 cm.

Trilateration is a process in which multiple receivers use the time of arrival signal from a tag to determine its position. This suffers from similar problems to other RFID methods in that it produces an area in which the tag could exist within - as opposed to its exact position. Combined with the need for 3 RFID readers, this approach fails to be accessible to the target audience.

The approach previously mentioned in Steve Hinske and Marc Langheinrich's paper, which utilised an antenna grid below the game board, seemed promising.

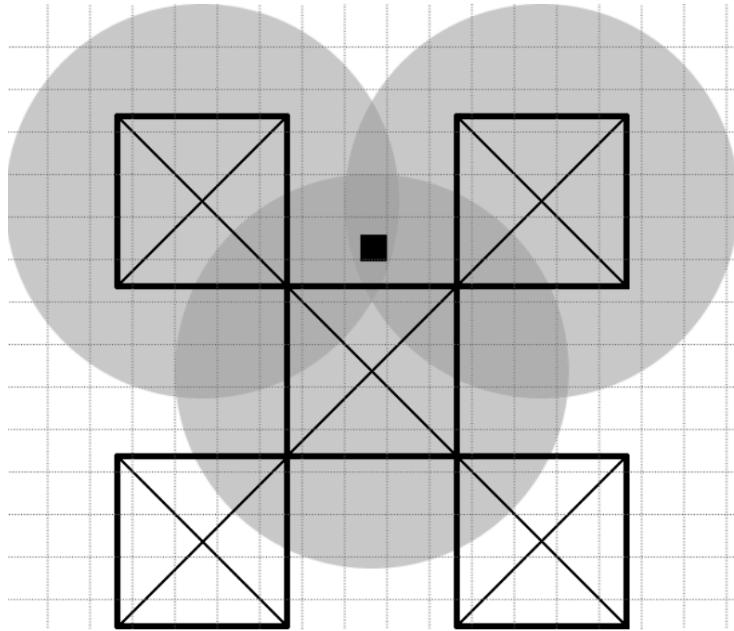


Figure 11: An example from Steve Hinske and Marc Langheinrich’s paper [11] describing their overlapping approach. The black square represents the RFID tag whilst the grey circles show the read area of each RFID reader. The intent is to use which readers are in range of the tag to determine an “area of uncertainty” of the tag’s location.

They were attempting to do this for Warhammer 40k, a game played on a much bigger board typically with larger models. As a result they were able to put RFID tags across larger models, such as vehicles, and utilise the known position of each tag relative to the model and the estimated position from the RFID readers to determine not only an accurate position but also orientation. Unfortunately for this project the size of the miniatures in Kill Team would require the use of one tag or two in very close proximity.

A potentially valid method utilising RFID could be to use the approach outlined by Zhongqin Wang, Ning Ye, Reza Malekian, Fu Xiao, and Ruchuan Wan in their paper [19] using first-order taylor series approximation to achieve mm level accuracy called TrackT. However, this approach is highly complex and can only track a small amount of tags at a time.

4.2.2. Machine / Deep Learning

Modern object tracking systems are often based on a machine learning or deep learning approach. In this case a classifier model would be used to identify each unique miniature on a team, and then subsequently locate it within the game board. The biggest drawback to this approach is the amount of training data needed for each class in a classifier. According to Darknet’s documentation (a framework for neural-networks such as YOLO) [20] each class should have 2000 training images.

Since each user will use their own miniatures, posed and painted in their own way, they would have to create their own dataset for their miniatures and train the model on them each time. As a user’s miniatures are likely to follow a similar paint scheme, ML classification could potentially struggle to identify between miniatures from top down and far away if not enough training data is supplied.

4.2.3. ARKit

Apple’s ARKit supports the scanning and detection of 3D objects [21] by finding 3D spatial features on a target. Currently any phone running IOS 12 or above is capable of utilising the ARKit. In this system,

you could scan in your models, then use the ARKit to detect them from above. This information could then be conveyed to the main system. This could also allow for the system to be expanded in the future to use a side-on camera as opposed to just top down detection, allowing for verticality within other Kill Team game systems. Combined with certain Apple products having an inbuilt LIDAR sensor (such as the iPhone 12+ Pro and iPad Pro 2020 - 2022), which further enhances the ARKit's detection, this could be a viable approach.

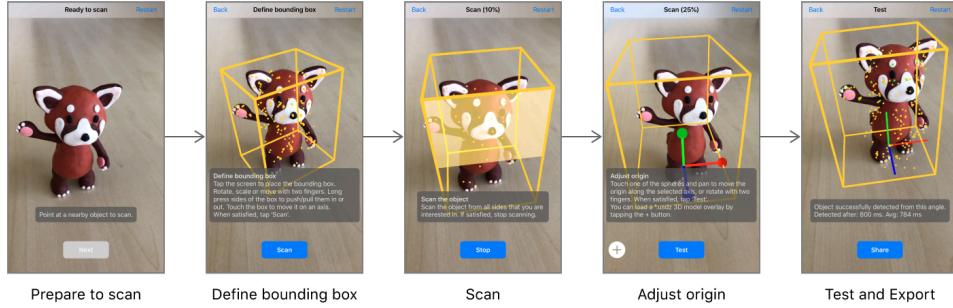


Figure 12: Registering an an object in ARKit [21]

After some testing utilising an iPad Pro with a LIDAR scanner, some drawbacks were found with this approach. The object detection required you to be too close to the models to detect them, and when being removed and re-added ARKit either took a long time to relocate a model or failed to do so at all.

As a result I won't be using this approach for the project. Although, for future work, this could be an interesting option to explore allowing you to implement AR features such as highlighting models on your phone or tablet or displaying information above the model.

4.2.4. Computer Vision

After considering all the options above, the best method found was using computer vision techniques.

One technique is blob detection. A blob is defined as a group of bright or dark pixels contrasting against a background. In this case the miniatures would be the blobs.

Utilising blob analysis would be good for locating the general position of miniatures (as they would very clearly stick out from the background) but getting their exact centre may then prove difficult (which would be needed for calculating movement and line of sight). Combined with the addition of terrain adding clutter to the image and the miniatures potentially being any colour, this approach could prove difficult to implement in the given time frame.

The computer vision method explored the most was utilising coloured tags to identify miniatures.

The use of coloured tags would also mean there would be access to specific measurements allowing for an accurate location to be discerned. However, some external modification to the game board or miniatures would need to be made. The challenge here is finding a way to do this without obstructing the flow of the game or the models themselves.

4.3. Fiducial Markers and Tag Detection

A promising approach appears to be using fiducial markers. Aruco tags [22] are a type of fiducial marker commonly used in robotics and computer vision. These tags are black and white squares with unique patterns in the centre. The tags are designed to be easily segmentable from the background and easily identifiable. The tags allow for pose estimation, which is the process of determining the rotation and translation of the marker in relation to the camera.



Figure 13: An example of aruco tags [22].

5. Project Description

This project aims to create a system that can track tabletop miniatures, in a game of Kill Team Gallowdark, using only materials easily accessible to the average miniature war-gamer. Then, utilise this system to create a digital representation of the game board. This digital representation will then be used to implement a few select game rules to demonstrate the system's capabilities and show that the tracking system provides the necessary information to process said rules.

The project can be broken down into two main goals.

1. Detection of the models and terrain to create a virtual representation of the game board.
 - a. The position of the miniature must be tracked accurately.
 - b. Be able to identify between different miniatures.
 - c. Aim to be non-invasive to the miniatures.
 - d. Be ambivalent to a model's shape and colour.
 - e. Be able to complete this task whilst being accessible to the average miniature war-gamer. Not utilising any equipment the average wargamer would not have access to in their own home.
2. Implementing the game logic in the virtual board to guide players through the game.
 - a. Allow users to select a model and which action they wish to preview.
 - b. Calculate the distance a model can move and display this on the virtual board.
 - a. Account for terrain that blocks movement.
 - c. Calculate the line of sight between the selected model and opposing models then display this on the virtual board.
 - a. Account for terrain that blocks line of sight.
 - b. Display the reasoning behind the line of sight calculation.
 - c. Display whether the target is obscured or in cover.
 - d. Display information about the selected model's odds to hit a target.

The methodology chosen to achieve these goals must meet the takeaways from the project background section. These requirements are not included here as they are more implementation specific instead of outlining the functionality of the project.

6. Methodology and Design

6.1. Tracking Methodology

We settled on using a computer vision and tag based approach to detect the models and terrain. This would utilise placing a camera above the centre of the game board on an adjustable stand and using tags to locate the models and terrain.

This approach was chosen for a few key reasons. Firstly, it is the most accessible option to the target audience. The only components needed are a camera, a computer, a printer and some method to position a camera above a board. Unlike RFID, IR or custom table methods, the average war-gamer would have access to these components. One intention for this project is to prove this is a viable method utilising only a phone camera and a stand.

Secondly, it is the most flexible option. A tag based system does not care about the shape or colour of the models, only the tags. This requirement meant that machine learning approaches would not be as viable as the potential variation in models used is too high. If we stuck to only supporting specific, unmodified models this may potentially work, but this would place a restriction on the target audience and in a subject such as miniature wargaming where customisation is so heavily valued, this would not be a good approach. As well as this given the unique paint schemes each player may use this would mean an ML model would need to be either trained specifically for each player's kill team or be able to use a model's silhouette to identify it. Both of these approaches either require users to make their own datasets and train the model themselves, a time consuming and technical task, or require a model that can identify a model from its silhouette, a complex task that would require a lot of training data and wouldn't be that accurate due to potential model customisations.

Building on this, maintaining a unique identification for an object is a key requirement but is difficult to implement using a machine learning approach. Kill Teams can utilise multiple of the same type of operative, as a result the system would need to be able to track each "unique" operative despite having the same appearance. This becomes problematic when occlusion is introduced. If we were to move an operative and at the same time cover or move another operative of the same type, it would be difficult to differentiate which one was which. A potential solution would be to utilise similar technology to facial recognition models, particularly looking at research into identifying the differences between identical twins.

In contrast, a tag based system is ambivalent to the appearance of the model. This means the system is functional with any model, regardless of customisation which meets one of the main requirements of the project.

Finally, a tag based system will assist with accuracy. Getting the location and rotation of an aurco tag using pose estimation is a well documented process. This methodology can be applied to a tag design that will function well with miniatures and terrain.

6.1.1. Camera Setup

A camera will be placed above the centre of the game board looking downwards providing a view of the entire game board. The image will be distorted by two factors: Camera distortion and image distortion.

Camera distortion is caused by different camera designs having slightly different distortion in the images they produce due to the different lenses. This can be corrected by using a camera calibration method to produce a camera matrix and distortion coefficients which can be used to undistort resultant

images. Radial distortions result in the image having straight lines appear curved. Tangential distortions result in the image appearing tilted along an axis. OpenCV provides a simple method to calibrate a camera using a chessboard pattern. This produces the camera matrix and distortion coefficients which can be saved and used to un-distort images. It is important to note that each different camera will have a different distortion so will need to be calibrated separately.

Image distortion is caused by the location of the camera relative to the board. To correct this we can perform perspective correction. Taking a top down view of the board will not guarantee that the board will be rectangular in the image. For example, it is likely for the camera to not be perfectly level resulting in parts of the board appearing closer or longer than they actually are. To remedy this we can identify the four corners of the board and perform a perspective correction algorithm to produce a flattened version of the board. At the same time this will give us the boundaries of the board and crop the image.

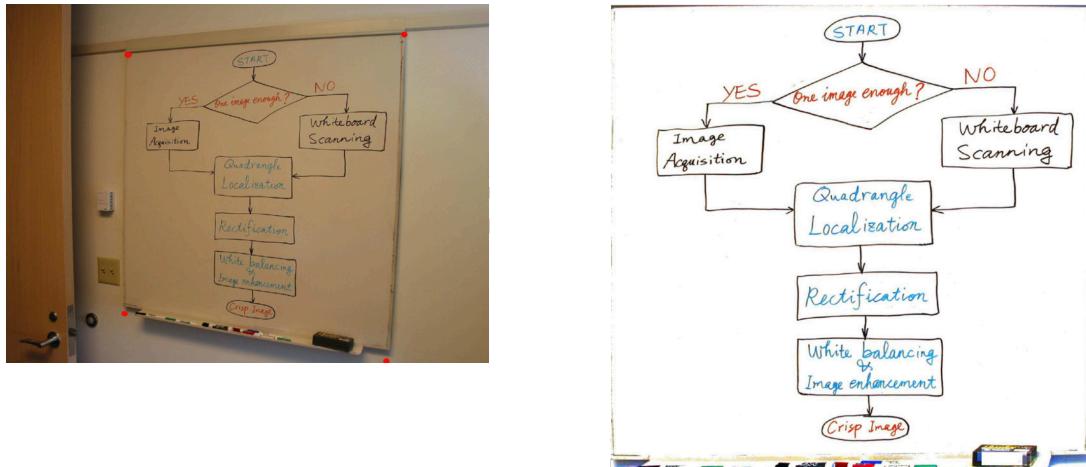


Figure 14: An example of perspective correction in Zhengyou Zhang and Li-Wei He [23]

6.1.2. Terrain Detection

Terrain detection utilises a straightforward approach. As the shape of terrain is limited in Kill Team Gallowdark to different constructions of pillars and walls. We can define a 2D model in the program to draw the terrain on the board. This means to detect terrain we only need to find a reference point to draw the model from, as opposed to needing to find the exact shape of the terrain.

This is achieved using aruco tags. As previously mentioned these allow us to perform pose estimation to find the location and rotation of the tag relative to the camera. Each type of terrain is defined by a unique aruco tag in a range. For example, a pillar with one wall can sit in the id range 1-10, a pillar with two walls in the id range 11-20 etc. Once we have identified the rotation, translation and scale of the tag we can apply these transformations to the model to draw the terrain on the board at the correct location.

The terrain detection system returns a list of terrain objects with:

1. The position of the four tag corners
2. The rotation of the tag as an angle
3. The id of the tag

6.1.3. Operative Detection and Identification

Operative detection is more complicated. Unlike terrain we can't place aruco tags on top of models as this is obstructive to the game and would appear significantly out of place. Instead we will design a tag to be placed around the base of the model.

The tag must be able to provide two pieces of information: the position of the operative and which operative it is.

6.1.3.1. Position

As the models are placed on top of circular bases the position of the operative can be defined by the centre point of the circle and the radius of its base. As the base size of a model is set, and the board size is also known, we do not have to worry about finding the radius of the base.

The requirements for the tag to provide the position are as such:

1. Be unobstructive to the model.
2. Be easily producible.
3. Be easily findable by the camera.
4. Provide the centre point of the operatives base.
5. Achieve all of the above whilst also being able to be occluded by both terrain and the model itself.
 - a. Ideally the centre point should be able to be found even with a quarter of the tag being visible.

These requirements resulted in this tag design:

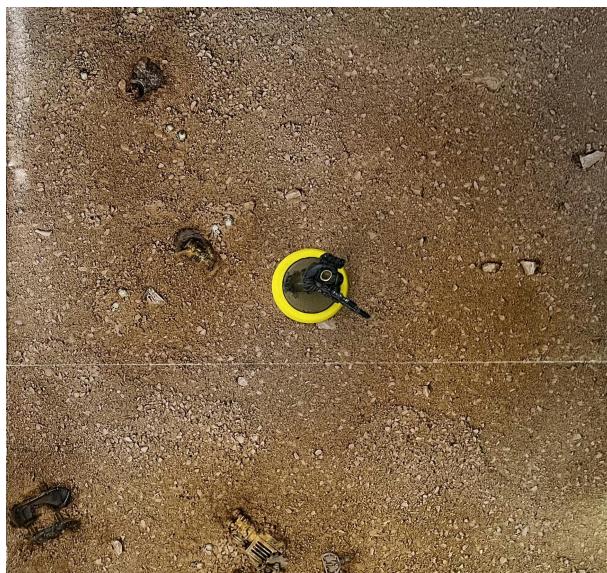


Figure 15: An example of the basic contrasting rim design.

To achieve this a high contrast rim will be placed around the base of the model. In this implementation we have chosen to use a bright yellow⁶.

Utilising Hough Circle transforms [24] in openCV we can easily find the centre point of a provided circle. The “completeness” of the curve required to be determined as a full circle can be easily adjusted to allow for occlusion.

⁶Although this could be changed to any colour which strongly contrasts the game board provided the correct thresholds were provided.



Figure 16: Circle detection on a game board with example terrain.

In extreme cases where the rim is heavily blocked by terrain the system will be unable to locate the model. However due to the nature of this project being aimed at following the Kill Team Gallowdark rule set, the terrain pieces used are placed along a grid. This prevents a situation where a terrain piece is very close to the edge of a board with a miniature placed behind resulting in the miniature being blocked from view.

6.1.3.2. Identification

The tag must also be able to provide a unique identifier for each model. The virtual game board will know which tag ID corresponds to which operative.

The method of identification on the tag must need to be functional whilst being occluded by terrain and the model itself. Ideally the tag should be able to be identified even if only a quarter of the tag is visible.

Kill Teams are typically made up of 7 - 14 operatives. This means our tag must be able to represent at least 28 different options.

To do this the following design was chosen:

It is important to note that the encoding is read right to left and uses little-endian⁷ encoding. When reading clockwise the first bit we encounter is the smallest. When reading anticlockwise the first bit we encounter is the largest.

⁷This is a little bit unconventional. Whilst this would appear to be big-endian, unlike binary, the encoding starts at the left and reads right so we encounter the smallest bit first, making this little endian.

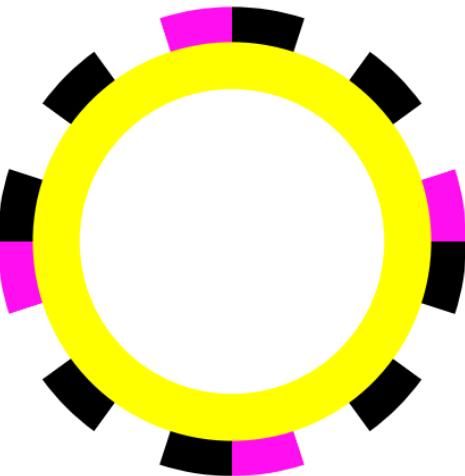


Figure 17: A tag representing the ID: 10. The outer ring is split into quarters with five sections within. The first section is a high contrast colour to indicate the start of the encoding. The four proceeding sections are split into black and white to represent the binary ID of the model. This pattern then repeats around the rim of the tag so only a portion of the tag needs to be visible to determine the ID.

Using 4 bits for identification allows for 16 different options. The first high contrast section colour can be changed to represent which team the model is on. This allows for 32 different options, more than enough for most Kill Team games.

The tag can be scaled to fit the size of the base of the model which is placed over the white circle. For this implementation we will use paper printed tags. This is a simple and cheap solution that is easy to get consistent colouring with. A tag of a similar design could be 3D printed with indents to paint in the colours.

Having an outer rim provides two benefits. Firstly, it allows for the identification bits to be placed further away from the model, making it less likely to be obscured. Secondly, it acts as a barrier to prevent the yellow rims from touching each other. This is important as Hough Circle detection can easily mistake two close but separate circles as one, larger circle or multiple smaller circles.

The high contrasting starting bit (shown here in magenta) is used to determine the starting point of the encoding. From here the system can then read clockwise and anti-clockwise to determine the binary ID.

6.2. Game Board Representation

The interface is created using Pygame. This is a simple to use library that allows for the creation of 2D games. It is relatively lightweight and quick to develop with.

The digital game board represents a 22" x 15" board. Which is half the size of a standard board.

6.2.1. Terrain and Operative Placement

The size of the image will likely be different to the size of the digital game board. Due to this we need to scale the terrain and operative positions, provided by the tracking system, to match the scale of the game board.

Operative positions are provided as a list of operative objects containing: the centre point and the ID.

The centre point of the circle is scaled and translated to match the game board scale. As the base size of the model is known, we draw a circle of the correct size at the scaled centre point. This keeps our digital representation accurate to the rules representation.

A list of operative ID's in use is stored. When an operative is detected, the ID is checked to see if it is in use. If it is in use, we update the position of that ID. If it is not in use, we ignore the detection as it was likely an error.

Terrain positions are provided as a list of terrain objects containing: the four corners of the tag, the rotation around the z axis relative to the camera and the id.

To draw the terrain we apply a scale to the model to match the game board scale. We then apply the rotation to the model before finally finding the translation between the tag corners and the corresponding points on the model. We then apply this translation to all the vertices of the model to then draw the terrain.

6.2.2. GUI design

The GUI is made up of a main game board window and the surrounding area. The surrounding area is used to display information not directly affecting the game board. Such as an explanation of what the board is representing, the controls for the system and a button to remove an operative from play.

The game board will display several things:

1. The operatives on the board, with their ID's and team colours.
2. Terrain coloured to represent whether it is light or heavy terrain.
3. Highlight the selected operative.
4. Once an operative is selected.
 - a. The line of sight status of the opposing operatives.
 - b. Opposing operatives who are in cover or obscured will be displayed in a different colour.
 - c. The firing cones of the selected operative to these obscured / covered operatives are also displayed.
 - d. Inside the firing cones the points providing cover and obscurement are also displayed.

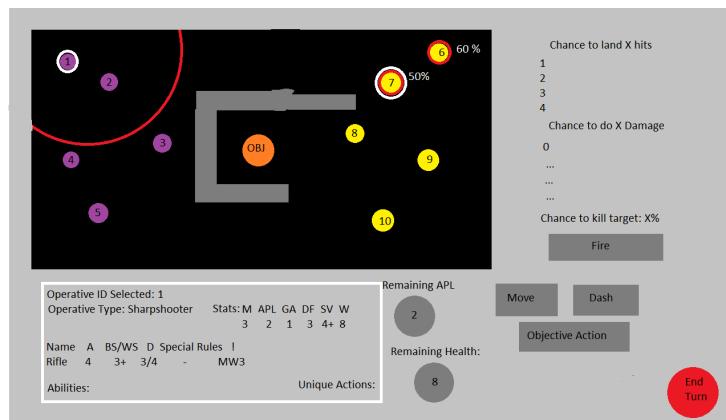


Figure 18: An example of what the GUI could look like⁸.

6.2.3. Line of Sight

In this implementation we handle the two complex parts of line of sight - obscured and in cover.

⁸The final GUI can be seen in the video and evaluation section.

6.2.3.1. Cover Lines

Both of these line of sight rules are based on the existence of terrain between the two models. The first thing we need to do is determine the sightline of a chosen operative to the opposing operatives. This is represented by two cover lines from a point on the attackers base to the extremes of the defenders base.

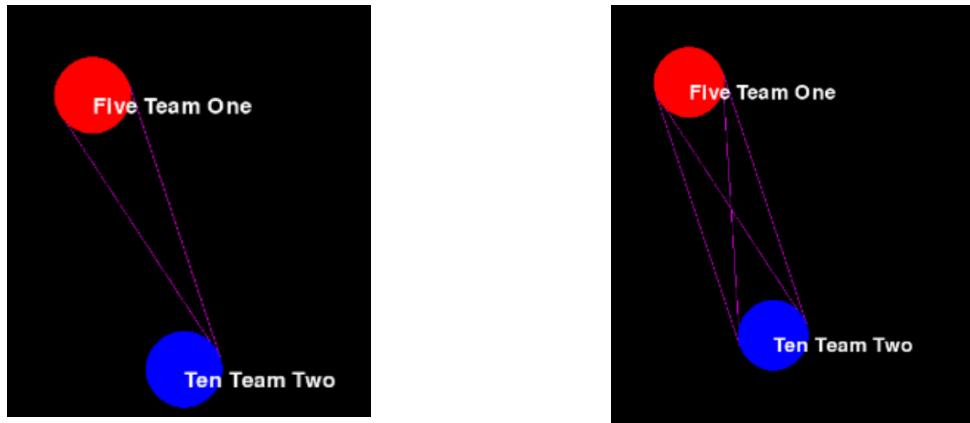


Figure 19: An example of cover lines being drawn. The left image shows the attacker (in blue) drawing a firing cone from one point on its base, to the extremes of the defender's base. The right image shows how this has been done from both sides of the attacker to get the two widest possible valid firing cones.

6.2.3.2. Obscuring and Covering Terrain

Once we have the firing cones we can determine whether an operative is obscured or in cover. This is done by producing a list of terrain lines that fall within the firing cones. We then check if any of the points on the terrain lines satisfy the distance requirements from the attacker and / or the defender. If they do, the obscuring points / cover points are drawn onto the board to provide an explanation to the user.

7. Implementation

This project will use openCV [25] to process the video feed from the camera. This library contains great support for computer vision tasks and is available for both Python and C++. An alternative image processing suite would be MATLAB but this is not directly compatible with other languages so would require a lot of extra work to integrate with a GUI.

Python was chosen as the language for this project due to both prior knowledge and Python's loose type system contributing to the ability to produce quick working prototypes separate from the main project. C++ on the other hand would produce a more robust final product but would take longer to develop. As this project aims to serve as a proof of concept for the idea, Python fits the role better.

7.1. Model Tracking

7.1.1. Design

A full sized kill team board is 22" x 30". This is too large for a single camera to capture the entire board and still be able to see miniatures behind terrain.



Figure 21: The board from up close.



Figure 20: An example of parallax.

Figure 22: The board from afar.

As seen in Figure 20, to get an image of the entire board, the camera needs to be $>1\text{m}$ away. Behind the wall, a yellow rim representing the model base is present.

This introduces three problems:

1. Having an arm long enough to hold the camera this high above the board is impractical and sometimes impossible given a small room.
2. The quality of the board in the image will decrease.
3. The colour of the board changes, as is visible in Figure 20, this would make our colour thresholding less effective and require more complex solutions to colour correct the image.
4. The image distortion will have a greater effect on the image. A small deviation from being level will have a greater effect on the image the further away the camera is.

The solution to this is to utilise two cameras.

One camera would cover the left half of the board and the other the right. This would allow for each camera to be closer to the board, whilst minimising the parallax effect. The two images can then be stitched together to create a single image of the entire board.

Whilst this was in the original plan, the project ended up focusing on creating a functional system for a half sized board so that, if there was time, a two camera solution could be implemented.

The camera design we ended up with was a single camera on a cheap phone holder stand pointed straight down at the centre of the board.

7.1.2. Video feed

To make the system as accessible as possible, we want to use a phone camera as the video input device. This would allow for the system to be used by the average war-gamer without needing to purchase any additional equipment.

To do this, we used a program called DroidCam [26]. This allows for a phone to connect either via USB or over wifi to a computer on the same network. The phone then acts as a webcam. The downside here is that droidcam is limited to 480p in the free version. Although, for the paid version the quality is increased to 1080p or even 4k if utilising the OBS plugin which you can then use to produce a virtual webcam.

For this use case, 1080p is more than enough.

This project was developed on an Arch linux with the 6.8.2 kernel. The phone used was an iPhone 13. As a result, when loading video input the system will use the linux method of video input. The

camera is represented as a file found in `/dev/videoX` where X is the number of the camera. 0 is usually the built-in camera and 2 is our external camera, though this can change depending on whether the external camera was connected on startup. This is easily changeable in the code in the *Camera* object to instead take an int instead of a “`/dev/videoX`” string for use on a windows system.

Getting video input from droidcam had two main issues.

The first was that the Arch package was broken. DroidCam makes use of the `v4l2loopback` kernel module to create a virtual webcam. As the video is not being produced by a physical card a virtual device has to be used. When kernel 6.8 was released `v4l2loopback` was broken. This left a few options to fix the issue. Either downgrade the kernel or compile the module from source with a community fix applied. Whilst the fix has been applied to the main branch it has not yet been released. Neither of these options were ideal.

Upon further examination it would appear that the default version for Arch is: `v4l2loopback-dkms 0.13.1-1` whereas droidcam makes use of a slightly different version: `v4l2loopback-dc-dkms 1:2.1.2-1`. According to the github page for DroidCam ,their version of `v4l2loopback-dc-dkms` has not been updated since 26/03/24. Arch kernel 6.8 was released on 29/03/24. It would also appear that DroidCam uses its own branch of `v4l2loopback` as indicated by the “dc” in the package name.

One user suggested installing the default branch of `v4l2loopback` instead. This still produced the same error, however after some further research it was found that running `sudo modprobe v4l2loopback` would fix the issue as the module was not being loaded on startup.

Once this was complete, the video feed was outputting in 1080p with low latency over wireless connection on Eduroam.

7.1.3. Calibration

To solve the camera distortion problem we need to calibrate the camera. This involves applying a camera calibration matrix to resultant images to account for the intrinsic properties of that specific camera design. There are two potential methods we could use. Most modern cameras have publically available lens profiles that can be used to correct the distortion [27]. Although these are aimed at use for photography or editing software, as a result they do not fit the format that OpenCV requires. The second method is to calibrate the camera ourselves.

OpenCV provides a simple method to calibrate a camera using a chessboard pattern [28].

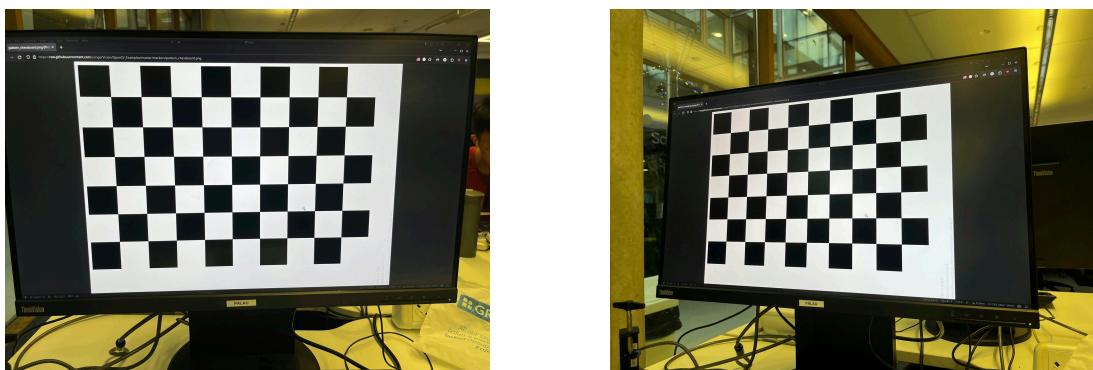


Figure 23: Some of the images used in calibration.

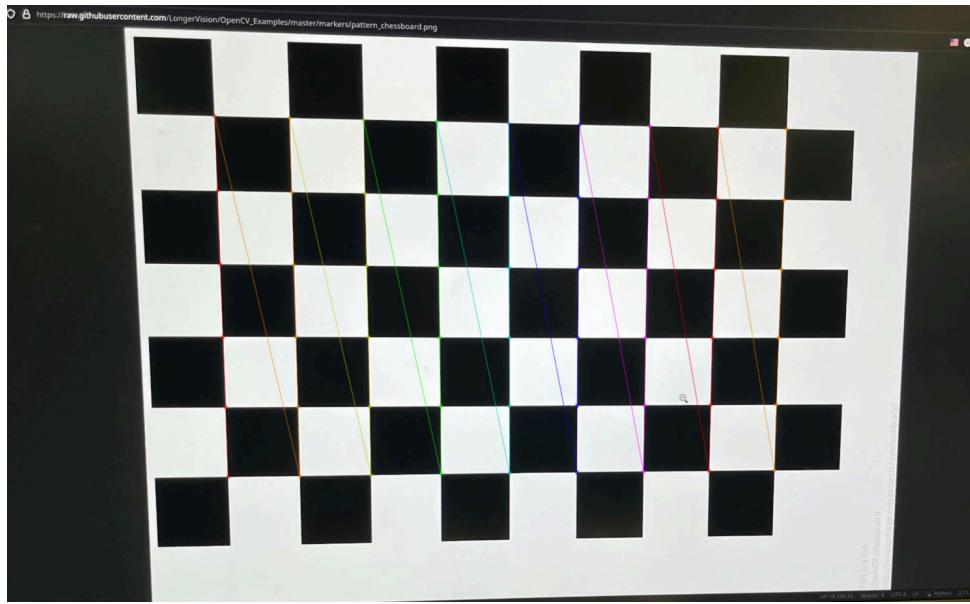


Figure 24: An example of the image points found in calibration.

A chessboard pattern is used as it is easy to find specific points of which the relative positions are known. In this case, the corners between black and white squares are used. As we know the position of these points in the real world and also the points at the image, we can use this to calculate the camera matrix and distortion coefficients.

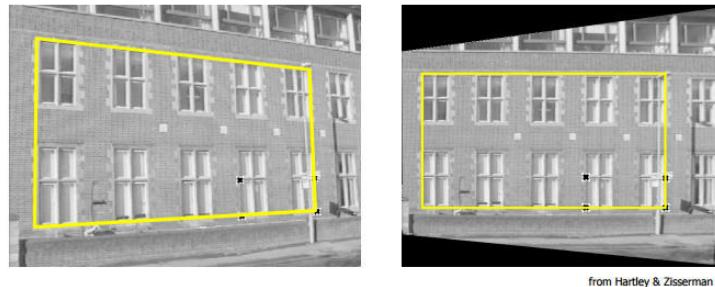
In the case of the camera matrix the points are used to calculate the focal length and optical centres of the camera. These values are then stored in a camera matrix.

The matrix and distortion coefficients are used later as they are required for the perspective-n-point pose computation utilised in the aruco tag detection.

Getting the calibration working took longer than expected. It is suggested to use around 15 images for calibration and when calibrating, each image was taking 3 - 5 minutes each and then not returning a calibration matrix. It was later found that this was due to the length and width of the chessboard pattern being passed as the wrong size. This was fixed and calibration now took <5 seconds in total. The length of calibration was due to OpenCV attempting to find a correctly size chessboard that was not present in the image.

7.1.4. Homography

Solving image distortion requires the use of homography. We want an image of the game board to be from a perfect top down view. But as the camera is not directly above the centre of the board, along with the camera not being perfectly level and distortion from the distance of the camera to the board, the board in the image will appear to be distorted. To correct this we need to find the transformation from the plane in the image to the plane as it should appear. This is known as performing a perspective transform, or perspective correction.



from Hartley & Zisserman

Figure 25: An example of perspective removal / correction from the OpenCV documentation [29].

To perform this transform to our game board we first need to find the four corners of the board. Ideally this would be done either through the use of aruco tags or through segmenting the board from the background. Board detection will be covered in section 7.1.5.

From these four corners we can form a trapezoid. Using this trapezoid, we can then find the height and width of the final rectangle⁹. With the new height and width we can then find the perspective transform with OpenCV’s `getPerspectiveTransform` function that correctly maps the trapezoid to the rectangle with minimal distortion. We can then apply this transform using OpenCV’s `warpPerspective` function to get a top down view of the board. This can be seen in Figure 26 and Figure 27



Figure 26: An example of perspective correction.

⁹This will be the largest height and width of the rectangle



Figure 27: A more extreme example of perspective correction.

7.1.5. Board Detection

An attempt was made at segmenting the board from the background as can be seen in Figure 28.

This approach utilised canny edge detection to find the separation between the board and the background and then HoughLines to find lines in the image. We then attempt to find a large rectangle from the provided lines. The image also has to be scaled down significantly for this to run in a reasonable time. Alternative approaches were taken using both contour detection and harris corner detection. However, these were not successful.

Contour detection appeared promising though the detailing on the board meant contours were incomplete. Attempts to fill the contours were made but unsuccessful. Automatic board detection was abandoned in favour of manually selecting the corners of the board on startup. The easiest solution would be to use aruco tags, but due to time constraints this was not implemented.

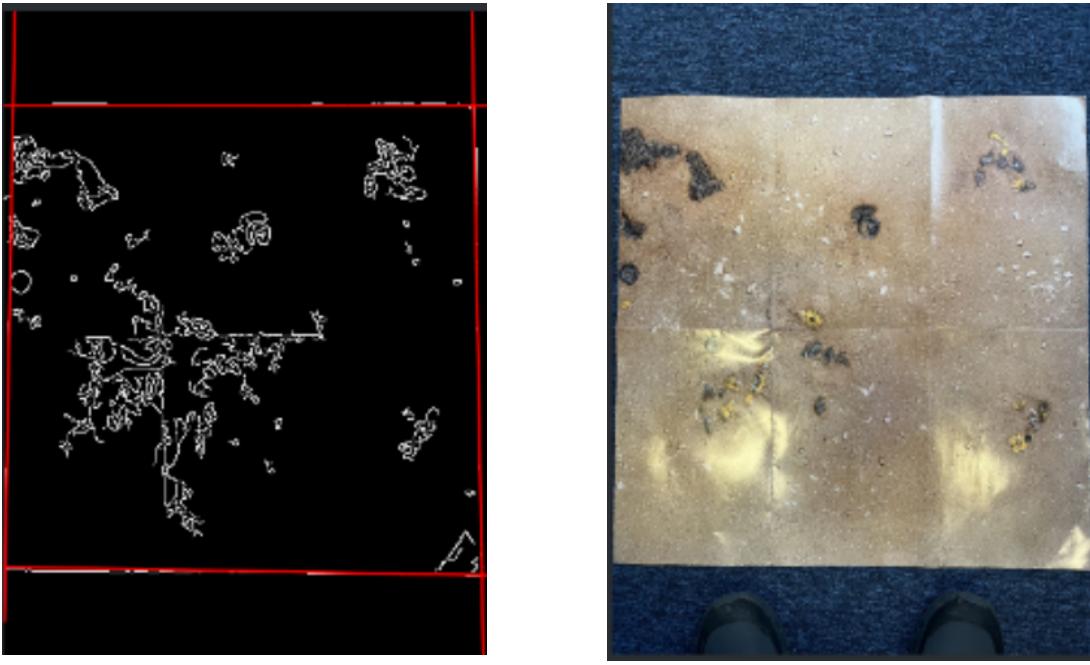


Figure 28: An attempt at segmenting the board from the background.

7.1.6. Tag Detection

Before we can attempt to locate the centre point of the circle we need to clear the image of noise. This is done by blurring the image¹⁰, converting to HSV and then applying a threshold to only show the yellow colour space in the image. This leaves us with a binary image with white pixels representing yellow and black for everything else.

From here, we perform edge detection on the image to find the edges of the circles in the image. This helps speed up the circle detection. This leaves us with a wireframe of the circles. We can then apply Hough Circle detection to the image to find the centre points of the circles and their radii. OpenCV provides a simple function to do this.

This leaves our processing pipeline as such:

1. Apply a gaussian blur to the image
2. Convert the image to HSV
3. Apply a threshold to only show the yellow colour space
4. Apply edge detection to the image
5. Apply Hough Circle detection to the image giving us a list of detected circle centre points and their radii.

It is important to note that Hough Circle transform can easily mistake two close but separate circles as one circle.

7.1.6.1. Hough Circles

OpenCV's implementation of Hough Circle transform makes up the backbone of the model tracking system, giving us the bases of each model.

¹⁰This is done to help reduce other noise from the image and soften edges.

The Hough Circle transform works by taking in a radius and drawing circles of that radius around each edge detected in the image. The result is an accumulator array where the highest values are where the most circles drawn by the transform intersect. This position is taken as the centre of the circle.

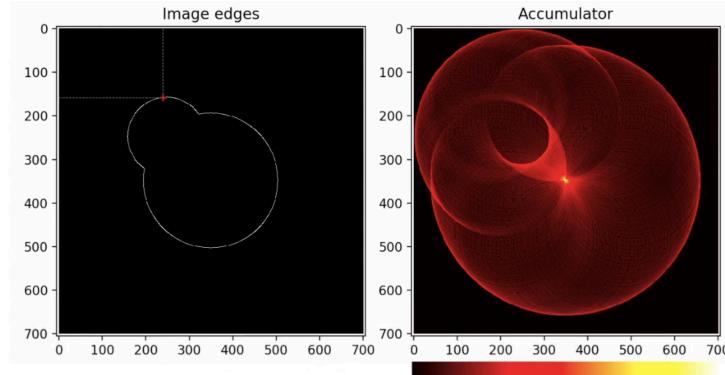


Figure 29: An example of Hough Circle detection and the accumulator produced. [30].

In OpenCV, this functionality is expanded to allow for both a range of radii, a minimum distance between the centres of circles and the minimum accumulator value needed to declare a circle centre. This minimum value allows for the transform to find incomplete circles. This is useful for our implementation as the yellow rims around the bases of the models are not always visible.

There is a problem with the use of Hough Circle detection. The perspective transform applied to the image can cause the circles to appear more like ellipses as they move away from the centre. This change is quite minor but in an ideal system we would utilise an ellipse or oval detection method instead. However, these implementations are not as readily available and would require more work to implement.

A large portion of this project was spent on finding and testing different approaches to detect the circles that was: robust, fast and allowed for identification to be done easily. Attempts were made using blob and contour detection, but these were not successful. Contour detection was not able to deal with occlusion well and blob detection struggled when the circles came close together.

7.1.7. Model Identification

The process to get the encoding is as follows:

1. Take the circle centres and radii from the Hough Circle detection.
2. Using the same transformed image as before, but unblurred and un-thresholded:
 - a. Convert the image to HSV.
 - b. Apply a gaussian blur.
 - c. Apply a threshold to only show the magenta colour space.
 - d. Perform an image dilation to ensure the magenta is a solid block.
 - e. Find the contours of the magenta.
 - f. Compute the centroids of the magenta.
 - g. Find the four closest centroids to each circle centre that are within the radius + a constant.

This will give us the starting positions of each visible encoding quarter for each circle.

From here we need to get the binary encoding.

The process to get the binary encoding is as follows:

1. For each circle:

- a. For each starting position provided (magenta centroid):
 - a. Rotate the coordinates of the magenta centroid around the circle centre to get the coordinates of each encoding bit.
 - b. This is done both clockwise and anti-clockwise.
 - c. Get the colour of the pixel at each encoding bit coordinate.
 - d. If the resultant colour is within the threshold values for white, the bit is 1. If it is within the threshold values for black, the bit is 0. Anything else returns a NaN value.

This will give us a circle centre, the associated magenta centroids and the clockwise and anticlockwise binary encoding associated with each magenta centroid.

Now that we have the colour values for each encoding bit in each quarter. We need to determine the final ID.

1. For each circle:
 - a. Reverse the anti-clockwise readings.
 - b. Group the associated encodings bits lists by their positions in the encoding to create 4 lists (position 0, position 1, etc)¹¹.
 - c. For each list:
 - a. The final encoding bit for that position is the majority bit present.
 - d. The final ID is the binary number created by the final encoding bits for each position.
 - e. Create an object containing the circle centre and the final ID.
 - f. Add this object to a list of all the found circles
 2. Return the list of all the found circles and their encodings.

7.1.8. Optimisations

At this point in the project the detection system was running, but it was exceptionally slow. Each detection (including image loading) was taking nearly 3 seconds as can be seen in Table 1. This was not acceptable for real time or even turn based system. The issue was found to be rotating the image to find the encoding bits. Instead of finding the coordinates of the encoding bits of each circle, the entire image was being rotated and the same coordinates check. This was a very expensive operation. Changing this to rotate the coordinates and not the image made significant time saves.

	Real (s)	User (s)	Sys (s)
Before	2.290	3.089	1.195
After	0.678	1.374	1.293

Table 1: The time taken for a single detection before and after rotation optimisation.

After chasing down several other instances of entire image operations, the time taken for a single detection was reduced to be mostly un-noticable. Providing the Python runtime would unfortunately not provide much information here as at this point, the program was just running the detection so a majority of the time was spent on overhead and image loading.

¹¹ for example 1 2 3 4 , 1 2 3 4, 1 2 2 3, 1 2 3 4 would become: [1,1,1,1] [2,2,2,2] [3,3,2,3] [4,4,3,4]

7.2. Terrain Detection

As mentioned previously, terrain detection is done using aruco tags. The aruco tags are placed on top of the pillars of the terrain so that the corners of the pillar align with the black corners of the tag.

OpenCV provides methods for performing pose estimation on aruco tags.



Figure 30: An example of pose estimation on an aruco tag.

Pose estimation will give us a matrix which tells us the rotation and translation of the tag in reference to the location of the camera.

The rotation of the aruco tag is provided in the form of a Rodrigues rotation vector. However, our model library requires an angle to perform a rotation. To convert between the two, we convert the rotation vector to a rotation matrix and then to Euler angles. From here we can extract the z-axis rotation (as we are working in a 2D plane from top down) and take the negative to get the correct angle for our model.

Terrain is slightly more complicated as it is defined as a series of 2D points to form a polygon. The terrain model is represented in mm sizes. The game board uses 3 pixels to represent 1mm. So we scale the terrain by 3x to match. We then rotate the terrain model to match the rotation of the tag. Finally we need to find the translation to move our model from model space into world space. This is done using the top left corner of the tag and the point which matches this on the model.

We scale the tag corner position into world space from the image space and then subtract the model corner position from the tag corner position to find the translation. From here, we can apply this translation to the terrain model to find the correct vertices to draw the terrain on the game board.

A terrain's ID is stored when it is detected originally. When that same ID is detected again, the position of the terrain is updated. If a new ID is detected, a new terrain object is created and added to the list of terrain objects.

Terrain detection caused more issues than expected. The main issue came from getting the rotation of the terrain. Converting from a rotation vector to Euler angles was a process that ended up taking a long time to get right. However, the main problem came from doing the pose estimation.

Semi recently, OpenCV was updated to include a new method for pose estimation, deprecating the old method. This new method was not as well documented as a result. The new method also added some levels of complexity. Previously, OpenCV had a dedicated function for pose estimation. Now, this functionality has been moved into solvePnP. This made the research done prior on implementing pose

estimation redundant. Eventually, a workaround was produced that allowed for the correct rotation to be found.

7.3. Game Board Representation

Originally, the plan was to use Qt for the GUI. However, this was abandoned for several reasons. Firstly, running Qt in a virtual environment on Arch caused issues for a while. Despite being installed, it simply would crash on startup. This was run into when trying to display images in OpenCV as it uses Qt for the GUI. This was solved by updating to Qt6 and then reverting to Qt5.

Secondly, when attempting to use Qt for the GUI, it would not display as desired. Qt is a very complex library and as a result getting simple GUI elements to display as desired was incredibly difficult. One of the main issues was trying to display the circles for the operatives. When placing a circle button in Qt it would place as desired, but adding multiple circles at the top of the screen would cause everything to shift. This is likely due to Qt being aimed at more generic user interfaces, rather than the specific use case of displaying a game board.

At the realisation that this project was closer to building a game than a general interface, the decision was made to switch to Pygame. Though a few other options were considered, such as Pyglet, Pygame was chosen due to its simplicity and abundance of documentation.

7.3.1. Linking of detected models and terrain to the virtual board

When a model is detected, its ID is checked against a list of IDs in use. If the ID is in use, the position of the operative in the list is updated. Before the position is updated, the coordinates from the image are translated to fit the board. This involves finding the scale between the virtual board and the board in the image as well as moving the coordinates to treat 0,0 as the top left corner of the virtual board and not the top left hand corner of the window.

If the ID is not in use, the tag is ignored. This is to prevent false positives from being added to the board.

As the game board is meant to be a strict representation of the rules, the base size of the model is defined within the program. This helps to ensure consistency between base sizes that would otherwise be difficult to maintain.

Operatives are also drawn with their team colours, state (concealed or engaged) and names.

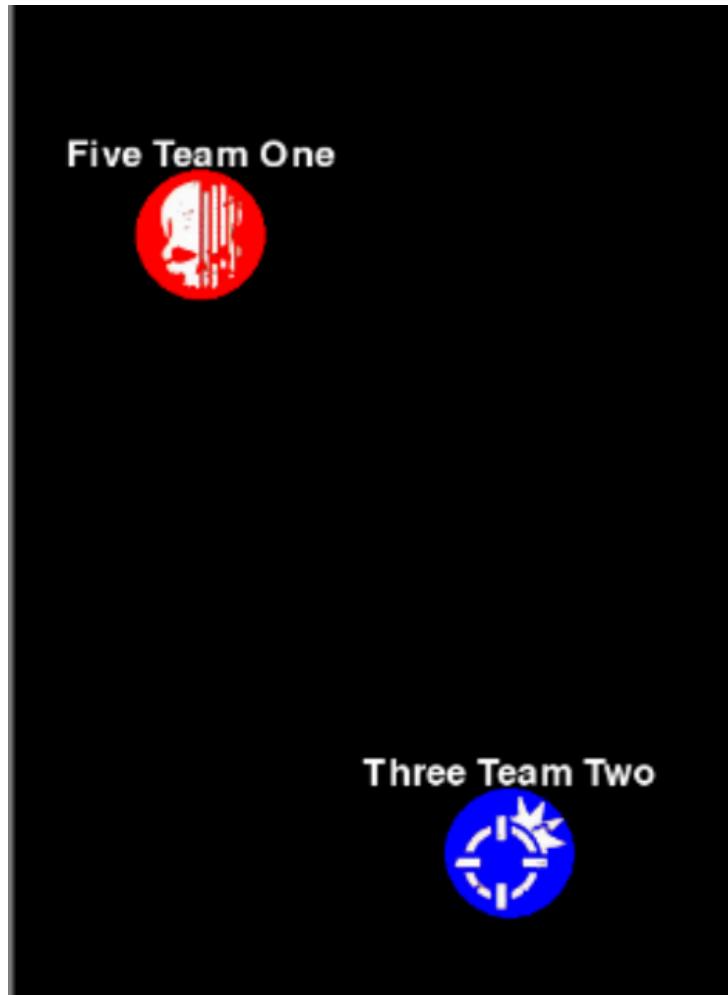


Figure 31: An example of the operatives and their state on the game board.

The relationship between models on the physical and virtual game board will be covered more in depth in section 8.1.

One issue that was encountered was that OpenCV uses y,x coordinates when indexing an image or getting the length and width. This caused some significant confusion when trying to find the required scaling between the image and the board.

Terrain functions in a similar way to operatives. Although, terrain does not need to be declared beforehand. When a new terrain ID is detected, a new terrain object is created and added to the list of terrain objects. If a terrain ID is detected that is already in use, the position and rotation of the terrain is updated.

The position of the terrain is based off of the top left hand corner of the aruco tag.

The terrain positioning ended up being broken for a while due to the terrain model being defined upside down. This meant that the top left hand corner of the tag was being matched to the bottom left hand corner of the terrain model.

7.3.2. Line of Sight

==W Firing Cones

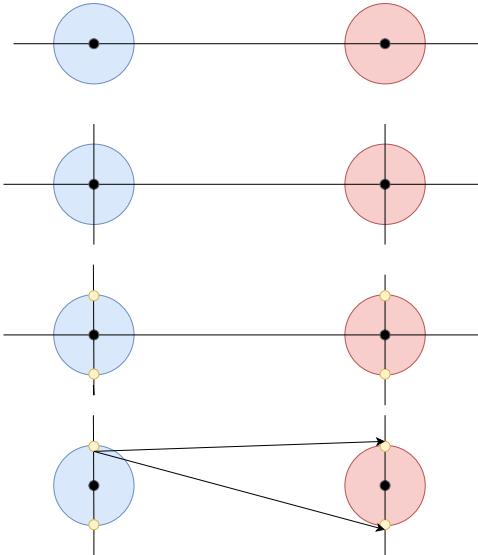


Figure 32: The process for finding the cover line positions.

As shown in Figure 32, the process for finding the cover lines is as follows:

1. Find the line equation between the two centre points of the bases.
2. Find the perpendicular line to this line through each of the centre points.
3. Find the intersection points of the perpendicular lines with the edges of the bases.

This gives us enough information to form the two firing cones we need.

Getting the firing cones is a simple process in theory but getting a functional implementation took significantly longer than expected.

Finding the line equation between the two centre points is a simple process. Getting the perpendicular gradient had some issues. Python had some problems with taking the reciprocal of a float.

Finding the intersection points on the bases proved problematic. Quite a few mistakes were made in the process of converting the algebra to Python code which took several days to notice.

A separate methodology was used based on a wolfram alpha solution. Although it was overlooked that the solution was specific to the circle being at 0,0 so this solution could not be used.

7.3.2.1. Terrain Within Firing Cones

Originally the plan was to use a ray-casting method to find the line of sight. This had two main problems. Firstly, it made getting the information required for obscured and in cover more difficult than it needed to be. As we are concerned with the distances of many points within the firing cones, a ray-casting implementation would need to be done past terrain after contact with the firing cone. Secondly, it would require a 2D array representation of the game board to be created. This would require some kind of rasterization of terrain and operatives to build. This would be very computationally intensive for a suboptimal solution.

The rules for obscured and in cover are very specific in their requirements as covered earlier. This allows us to exploit the specifics of the problem to simplify the process. As the line of sight rules are more complicated than simply whether an operative is visible to another operative, we need a unique solution to this problem.

We can break the problem down into several parts:

1. Building the firing cone.

2. Determining what terrain is within the firing cone
 - a. Creating a list of lines within the firing cone.
3. For each line finding the closest point to the operative.
4. Determining whether the point meets the obscured or in cover requirements.

Both obscured and in cover are determined by the distance between the attacker / defender and a point at which terrain is within the firing cone.

We can determine whether a terrain line is within the firing cone if it satisfies any combination of the following two conditions:

1. The start or end of the line falls within the firing cone.
2. The line intersects with the firing cone.

Using this we can rebuild a list of terrain lines that fall within the firing cones.

The lines are determined as follows:

1. If a line satisfies none of these then it can be ignored.
2. If it satisfies both points within the firing cone then we take both the start and end points.
3. If a line satisfies only one point within the firing cone then the line must also intersect with the firing cone. We then take the point within the firing cone and the point of intersection.
4. If a line intersects the firing cone at two points then we take both points of intersection.

This leaves us with a list of lines that fall within the firing cone.

Finding whether a point falls within the firing cone is done by calculating the barycentric coordinates of the point in relation to the triangle formed between the two points on the defender and the single point on the attacker. If alpha, beta, and gamma are all between 0 and 1, then that point falls within the triangle.

Finding whether a line intersects with the firing cone is a bit trickier. Finding the intersection of a line and another line is easy. However, finding the intersection between two line segments is more complicated. This is because the intersection point of two lines can be outside of the line segments. Whilst an algorithm existed to solve this, it had several bugs that needed to be fixed.

Now that we have a list of lines that fall within the firing cone we can find the closest point on each line to an operative. If the point meets the obscured or in cover requirements then we check if the requirements are also met for the attacker / defender. If they are, then we set the defender to be obscured or in cover.

Finding the closest point on a line to another point is a deceptively difficult problem. The solution ended up using some vector algebra to project the point onto the line using a pre-existing algorithm.

8. Evaluation

This section will discuss the results and limitations of the project. Section 7.1 will discuss how well the system performs in tracking models and terrain. Section 7.2 will discuss the overall outcome of the project in relation to the initial goals set out in section 4.

8.1. Tracking System

The first test involves comparing the position of the detected models on the board to the actual position of the models.

The second test involves checking how well the system can track multiple models at once.

As this is a live system, the tests will be performed in real time. This is difficult to demonstrate in a report so screenshots will be provided to show the results of the tests along with a description of observations. These observations can be better seen in the supplied demonstration video.

8.1.1. Accuracy

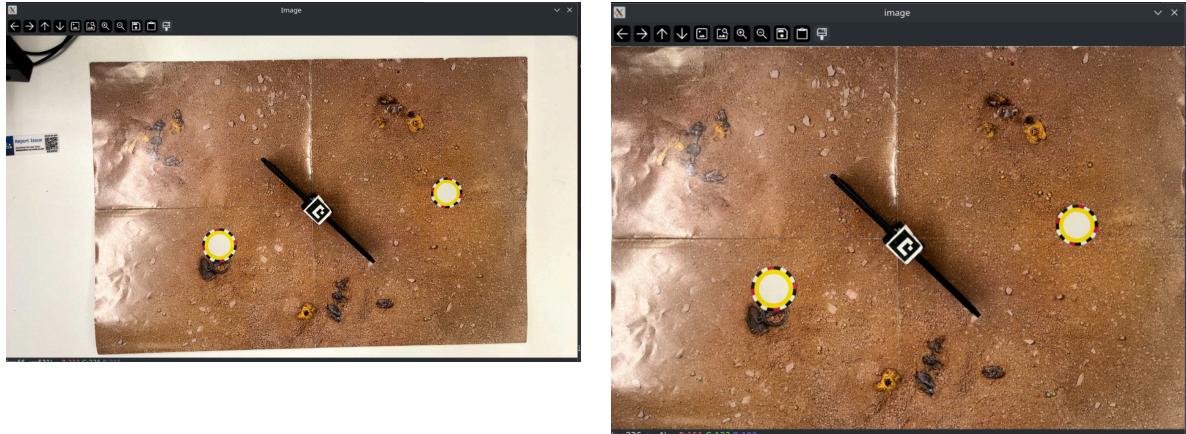


Figure 33: The actual board and the transformed board¹².

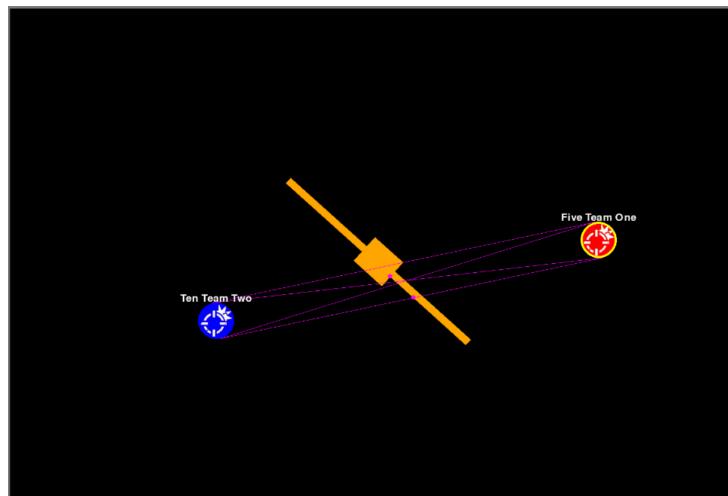


Figure 34: The digital game board of Figure 33

The tracking system provided the coordinates of (754,771) for “Ten Team Two” (pictured in blue) and (1642,581) for “Five Team One” (pictured in red) as seen in Figure 34. These values are including the translation to fit the Pygame window, which is 279 pixels in the x direction and 50 in the y. Each pixel representing 3mm. (0,0) is considered to be at the top left corner of the board with y increasing downwards.

The equation to convert from the Pygame coordinates to the real world is as follows:

Let a and b be the game board coordinates for x and y respectively. Let x and y be the actual, real world coordinates for x and y respectively.

$$x = \frac{a-279}{3}$$

$$y = \frac{b-50}{3}$$

¹²It is important to note that the binary decoding for these tests was reversed - this has since been fixed. But the encodings on the rings are swapped as they're symmetrical.

These values represent the position in mm from the top left corner of the board in the x and y directions respectively.

Operative Name	Pygame	game board	Real World	Actual	Difference
ID: Ten, Team Two (blue)	(754,771)	(475,721)	(158,240)	(160,235)	(2,5)
ID: Five, Team One (red)	(1642,581)	(1363,531)	(454,177)	(459,181)	(5,4)

Table 2: Results of the tracking system from Figure 33

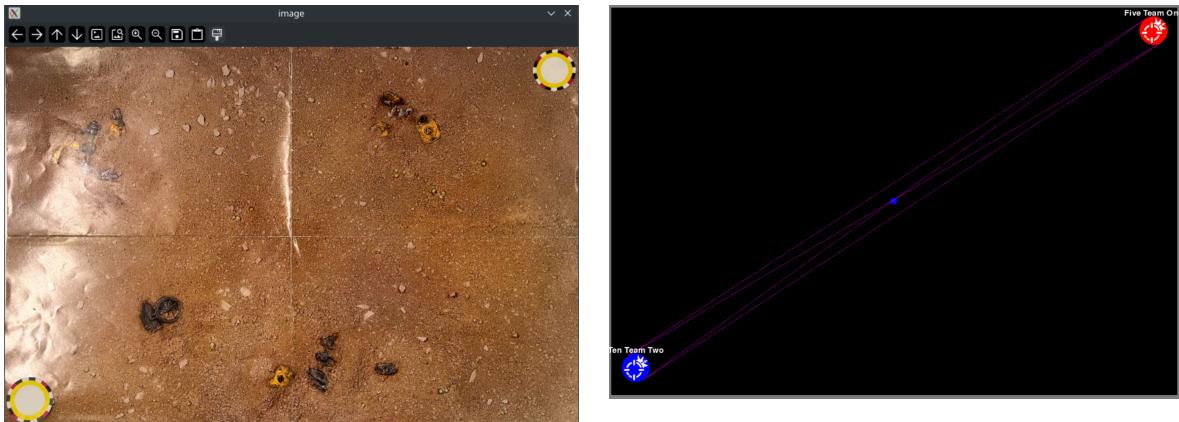


Figure 35: Tracking at the corners of the board to see how distance affects detection.

Operative Name	Pygame	game board	Real World	Actual	Difference
ID: Ten, Team Two (blue)	(352,1113)	(73,1063)	(24,354)	(27,349)	(3,5)
ID: Five, Team One (red)	(1886,118)	(1607,68)	(536,23)	(535,24)	(1,1)

Table 3: Results of the tracking system from Figure 35¹³

These results are very promising. The tracking system seems to be able to accurately determine the position of the models on the board down to 5mm. This is a very good result. The tags were marked on the centre of the base when taking the real world measurements. Although, these were not perfectly centred. A visualisation of the accuracy can be better seen in Figure 39.

The next test will look at how well the system can track operatives with the models present.

¹³The small blue circle in the right image is the board centre point.



Figure 36: Tracking with models present.

In Figure 36 the system correctly identified the red team model but failed to identify the blue team model. The tracking system determined that the blue team model was ID: 0 or ID: 4. This is likely due to the blue team model being partially cropped out of the image and containing black on the model. As a result when checking for the binary parts of the model are interfering.

However, as seen in Figure 37 the system was able to eventually identify the blue team model. This presents a limitation of the current system. As the system does not have a confidence value for the identification, it will simply return any valid identification it can find. This will result in problems when more models are present as the system will be likely to move a model into another when mis-identifications are present. A potential fix for this would be to utilise a system using hamming distance to better identify models with partial data. This will be expanded upon in Section 8.2. A temporary fix was applied that would prevent the system from placing a model in the same position as a model that was already present. Whilst this isn't a particularly robust solution, it does prevent the system from placing two models in the same position.

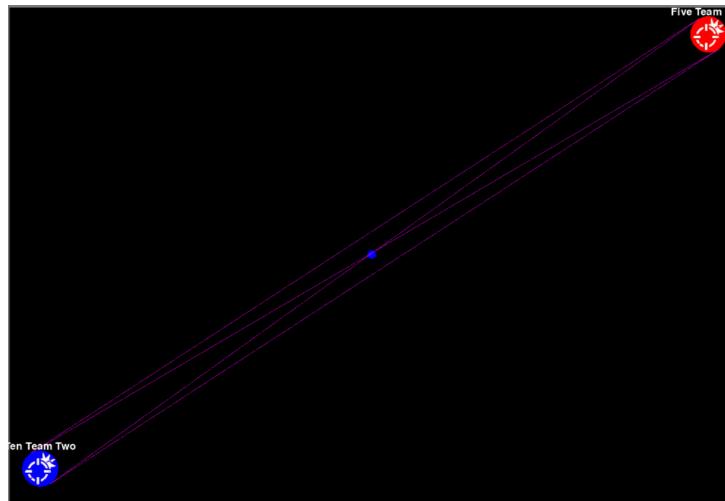


Figure 37: The same setup as Figure 36 but 5 frames later.

8.1.2. Multiple Tag Tracking

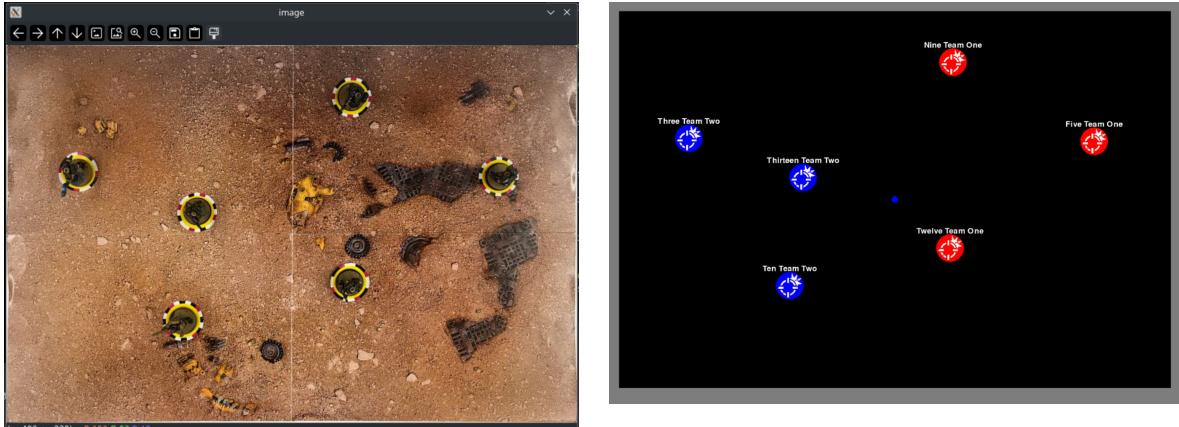


Figure 38: The game board with 6 models present.



Figure 39: The images in Figure 38 overlaid on each other.

As seen in Figure 39, the system was able to track all six models present. The system is able to track multiple tags, with the correct identification, at the same time with high levels of accuracy. Although, it would appear that as the tags move further up, the accuracy dwindles slightly, this can be seen in “Nine Team One” in Figure 39.

8.1.3. Tracking With Terrain



Figure 40: The game board with 6 models and terrain present.



Figure 41: The images in Figure 40 overlaid on each other.

As seen in Figure 41, the system is able to track both models and terrain simultaneously. However, the terrain appears to be slightly tilted to the right. This is likely due to the position of the aruco tag on the terrain being slightly offset. As well as this, you can see that the top left corner of the aruco tag and the top left hand corner of the terrain pillar are not well lined up. This will contribute to the mis-positioning of the terrain. Although in relation to the models and the centre of the board, the terrain is in a good position.



Figure 42: The game board with a model obscured by terrain.

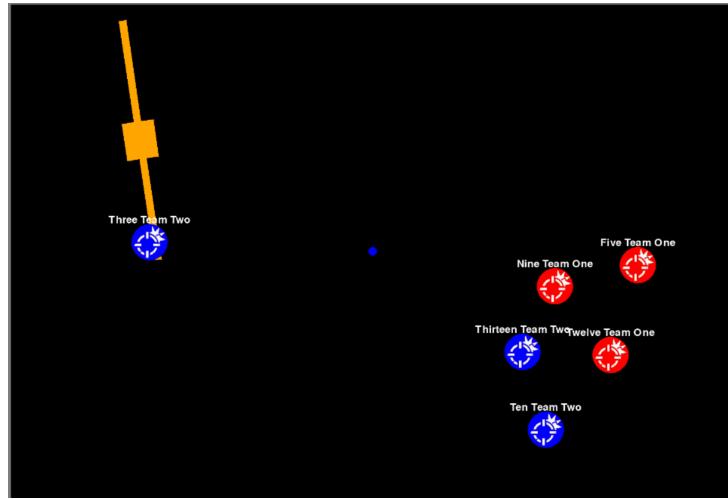


Figure 43: The digital game board with a model obscured by terrain.

In Figure 43 the system was unable to find the obscured model. The use of black terrain is likely causing the system to find 0s in the binary encoding, making identification difficult when combined with the blocking of the starting bits. However, as seen in Figure 44 when rotating the model the system is able to find a correct identification.

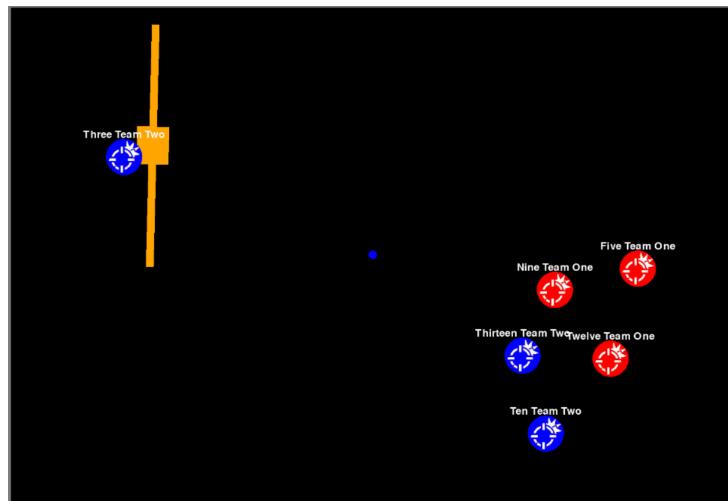


Figure 44: The digital game board after the model was rotated.

Another problem is present which is visible in Figure 44. The roll of the terrain is not being accounted for, resulting in an incorrect positioning. A potential fix for this will be discussed in section 8.2.



Figure 45: Figure 42 overlayed with Figure 43

8.1.4. Summary

Some testing was done to compare the usefulness of the visualisation of the game board. This involved playing a solo round of kill team with the system with four operatives. Overall, the system was able to track and identify miniatures to a satisfactory level. There were a few times where the system would take a while to find a model but this was usually fixable with some rotating and slight shifting. However, terrain ended up being problematic. The system could get the rough area of terrain but struggled to get the exact position, combined with the terrain shifting slightly whilst stationary, this made the experience less enjoyable. The terrain shifting is likely caused by a lack of distance checking when updating terrain. Whilst the terrain rotation must be greater than a constant value to be updated, this was overlooked for the position. Some different approaches to remedy this will be discussed in section 8.2.

The line of sight visualisation was satisfactory, displaying in cover or obscured to a degree that appeared realistic.

9. Summary and Reflections

9.1. Project Management

The first and second gannt charts show the original and interim gannt charts for the project¹⁴. The final gannt chart shows the updated gannt chart with the actual timeline of the project.

Previously, I stated that tackling the ring detection first was a good choice. This allowed for me to encounter larger issues and provide solutions to them whilst still in the research stages on the project. However, in hindsight tackling terrain would have been a better decision. As the terrain detection system was using pre-existing methods, it would have taught me about how fiducial marker systems work. This would have been useful when building the model detection system.

On the other hand the model detection system took a long time to get right. Leaving this until later would have likely caused significant stress as it is the main focus of the project. It was a part of the project which was very important to get right, so leaving plenty of time was the sensible choice to make.

¹⁴These have been submitted in the appendix.

Supervisor meetings were held every week to check for blockers or advice on reports. Uniquely, our supervisor meetings were conducted as a group with the other two dissertation students my supervisor had. As all three of us were working on tabletop game based projects we were all knowledgeable in the area. This meant we could provide feedback or ideas for each other's projects from a student perspective.

Previously I mentioned that gannt charts were unhelpful in managing the project and that I would instead take a kanban based approach from my previous successes in the methodology in the second year group project. In reality, past the first week this was left behind. Instead, I opted for leaving myself a comment block about what I had achieved that day and what I needed to do next. This came in useful when writing the report as I had a breakdown of what was completed and the difficulties faced that I would otherwise not have remembered. This was a much more effective method of tracking progress than the gantt chart.

I also previously mentioned that this project was taken as a 70 / 50 split. This made the second half of the project significantly more manageable. Although this project did make use of extension time, that time fell within the easter break. I find that development work is much more effective when done in long, uninterrupted sessions. being able to have only 50 credits of work to do with a long break at the end of the project was very helpful.

I achieved the majority of what I wanted to get done in this project with the only missing part being movement preview. This was not a priority as the technical interest of this project is the development of the tracking system.

9.2. Future Work and Reflections

Something I underestimated was the amount of time needed to research and determine which method should be used for model detection. A large amount of time was spent determining the viability of different approaches that may make the system more robust and easier to produce. As a result of this, development started later and took longer than was expected.

The biggest impact on this project was needing to get an extension on other coursework deadlines. This extension had a knock on effect of other courseworks which took priority over this project. The timings allotted in the Gantt chart did not take into account extensions being needed.

The reflection above was written at the time of interim report submission. It has been left in as I believe it is important to evaluate how I handled the project in the past compared to now.

I started development by attempting to produce a tag detection system for the operatives and then move onto terrain afterwards. In hindsight, it would have been better to start with terrain detection. I utilised aruco markers to find the position of terrain. The process of implementing aruco markers ended up teaching a lot about fiducial markers that I wish I had known before building the model tags. I'll cover the specifics of this when discussing upgrades to the detection system.

9.2.1. Order Tokens

Order tokens are a key part of the game that are not currently implemented. These are usually small pieces of orange triangular card pointing at an operative with an image denoting the state. A future system might be able to segment these tokens out and use the direction of the triangle to determine the state of an operative without needing manual input.

9.2.2. Odds to Hit

As Kill Team is mostly based on dice rolls modified by the statistics of the operatives, a future system could implement a way to calculate the odds of hitting a target. A problem with this would be needing to then store and update health values manually. This could be done by moving the dice roll calculation to being digital. Although, I don't think this is a good idea. An important part of tabletop boardgames is actually physically rolling the dice.

There are plenty of pre-existing systems to see the values of physical dice. These typically either utilise computer vision to detect the sides of the dice or use specific dice that are aware of their orientation. A "smart dice" system would prevent the need for manual interaction with the system. This would be a good place to start for a next addition to the system as it opens up a lot more game mechanics to be streamlines.

9.2.3. Increasing total number of operatives

The current system only supports 16 operatives total on a half-size board. The colours of the encodings can be changed to allow for more but this requires thresholding for different colour values to be re-done. A better approach may be to use more than 4 bits for the encoding and utilise hamming distance to deal with partial data

9.2.4. Server Client Architecture For Game Board and Detection

The current system is run on a single thread. This is fine as it currently is but there are a few problems. One such problem is that the GUI is running on the same thread as the detection meaning responsiveness is likely to be a problem on slower machines. A solution to this would be to have the detection system function as a server that can serve data to the game boards.

This also opens up the avenue for alternative detection systems to be used as a server based architecture would allow for the creation of an API that is agnostic to the type of detection used.

Alternatively this could be expanded to allow two players to play against each other on different game boards by transmitting their model positions to each other.

9.2.5. Solo Play

If the system were to receive more rules integrations it may be possible to make an AI that can play against a human. This would require a huge amount of work to implement the rules of the game but it would be very interesting to see.

9.2.6. Detection Upgrades

One of the biggest changes that could be made would be to make the operative tags utilise hamming distance. Hamming distance is a method that lets us maximise the difference between our marker patterns. If we design markers that are maximally different from each other then when errors are present we can use the information we do have to determine which marker is most likely to be the correct one. For example. if we're using a 4 bit encoding but we only have two tags 0,1,1,1 and 1,0,0,0 we can determine the ID of the tag based on only one bit of information. This method can be applied in a more complex way to develop a tagset. This is a system that is used by aruco tags.

Another problem is with the positioning of terrain. The terrain is only rotated in the Z axis from the aruco markers which is resulting in the terrain being slightly off. In future a different model library should be used to allow for more complex rotations to increase the accuracy of the terrain.

The same pose estimation approach should also likely be taken to the operative tags. This system would greatly benefit from a more robust method of moving models. Currently everything is done by x,y coordinate movement but it would be much more streamlined to use a proper matrix transformation system to move models. This would also make the system much more maintainable as the current system does a lot of calculations between different spaces all over the place and having a streamlined pipeline would make this much easier to manage.

9.3. LSEPI

The main LSEPI issue you could see here is copyright issues as Kill Team is a copyrighted entity owned by Games Workshop. The way this project will handle this is by having the system implement the Kill Team Lite rule set previously mentioned. This is publicly available published by Games Workshop. One issue is that different Kill Teams will have different and unique rules as well as their own statistics pages. These are copyrighted and won't be able to be directly implemented. As a result of this, this project will implement basic operatives with fake statistic pages based off of the publicly published Kill Team information. If this proves to be problematic then the game rules will be based on a very similar system Firefight [31]. This is published by One Page Rules, who produce free to use, public miniature war-game systems. Though as this is based off of the published, free to use Kill Team Lite rules, I don't expect this to be a problem.

9.3.1. Accessibility

A key focus of this project was on accessibility to the average-miniature wargamer. This led to a focus on not using specialist hardware. I believe the system achieved this goal with the main potential issues being technical setup and camera setup. Being a Python project, Python needs to be installed with the libraries. The process of doing this to a non-technical individual can appear quite daunting. The camera setup is also a potential issue. Droidcam ended up being a harder to use piece of software than I first anticipated. This could be a potential issue for the average user.

9.3.2. Open Source

I intend to open source the code for this project using a copyleft licence. When performing the background research for this project I found that there were very few closely linked projects. As a result of this I want the code and the takeaways I gathered from this project to be available to the public to assist anyone undertaking a similar project in the future.

10. Final Reflections

Overall, I believe this project has met the goals of digitising a physical game board. The accuracy of the model tracking is very good and the terrain tracking is satisfactory but needs some improvements. The line of sight system provides enough information to the user to know whether a model is obscured / in cover as well as what is causing the obstruction.

A problem I ran into quite frequently was the OpenCV documentation was quite lacking. Documentation for versions >4.7 tended to only support C++. OpenCV claims to have support for 5 different languages but only provides in-depth documentation on C++. From the work I have done, it appears that a lot of questions that get asked about OpenCV are asked in the context of Python. This is a problem I have run into in several open source projects. Whilst the functionality is amazing the documentation tends to go out of date quickly as contributions are made. This is a shame as the OpenCV documenta-

tion which is of good quality is incredibly well done. It explains the concepts at both a low and high level with examples. A big takeaway from this project will be in contributing to the documentation of open source projects I use if I can. OpenCV does have a documentation repository that I will likely contribute to in the future to provide more Python examples.

11. Bibliography

- [1] Games Workshop, “Games Workshop Investor Relations Statement.” Oct. 17, 2023. [Online]. Available: <https://investor.games-workshop.com/our-history>
- [2] Games Workshop, “Warhammer 40k Core Rules.” Jun. 02, 2022. [Online]. Available: <https://www.warhammer-community.com/wp-content/uploads/2023/06/dLZIlatQJ3qOkGP7.pdf>
- [3] Games Workshop, “Kill Team Lite Rules.” Aug. 16, 2022. [Online]. Available: <https://www.warhammer-community.com/wp-content/uploads/2022/08/ekD0GG2pTHlYba0G.pdf>
- [4] Warhammer Community, “Kill Team: Into the Dark – What’s in the Box?” Accessed: Dec. 31, 2023. [Online]. Available: <https://www.warhammer-community.com/2022/08/04/kill-team-into-the-dark-whats-in-the-box/>
- [5] Games Workshop, “Killzone Gallowdark.” Accessed: Dec. 12, 2023. [Online]. Available: <https://www.warhammer.com/en-GB/shop/killzone-gallowdark-2023>
- [6] Games Worksshop, “Kill Team: Kasrkin.” [Online]. Available: <https://www.warhammer.com/en-GB/shop/kill-team-kasrkin-2023?queryID=4bbc70cecccea9006fc261421f8bc787>
- [7] Warhammer Community, “Sneaking and Spotting Are Just as Important as Shooting and Stabbing in the New Kill Team .” [Online]. Available: <https://www.warhammer-community.com/2021/08/11/sneaking-and-spotting-are-just-as-important-as-shooting-and-stabbing-in-the-new-kill-team/>
- [8] Last Gameboard, “The Last Gameboard Kickstarter.” Accessed: Dec. 12, 2023. [Online]. Available: <https://www.kickstarter.com/projects/gameboard1/gameboard-1>
- [9] Teburu, “Teburu.” Accessed: Dec. 12, 2023. [Online]. Available: <https://teburu.net/#home>
- [10] “Teburu: a sneak peek.” Accessed: Dec. 13, 2023. [Online]. Available: <https://www.youtube.com/watch?v=5paNoKA4b5A>
- [11] Steve Hinske and Marc Langheinrich, “An RFID-based Infrastructure for Automatically Determining the Position and Orientation of Game Objects in Tabletop Games.” Jun. 05, 2007. [Online]. Available: <https://vs.inf.ethz.ch/publ/papers/hinske-pg07-rfidtabletop.pdf>
- [12] Whitney Babcock-McConnell, Michael Cole, Stephen Dewhurst, Bulut Karakaya, Dyala Kattan-Wright, and Maokai Xiao, “Surfacescapes.” Accessed: Dec. 13, 2023. [Online]. Available: <https://www.etc.cmu.edu/projects/surfacescapes/index.html>
- [13] Microsoft, “Microsoft Surface 1.0 documentation.” Accessed: Dec. 13, 2023. [Online]. Available: <https://social.technet.microsoft.com/wiki/contents/articles/8017.microsoft-surface-1-0-sp1-getting-started-guide-hardware-overview.aspx#Tabletop>
- [14] Abhay Buch, “With Surfacescapes, Dungeons & Dragons becomes high-tech.” Accessed: Dec. 13, 2023. [Online]. Available: <http://thetartan.org/2010/4/5/scitech/surfacescapes>
- [15] Foundry Gaming LLC, “Foundry VTT.” Accessed: Dec. 23, 2023. [Online]. Available: <https://foundryvtt.com/>

- [16] Material Foundry, “Material Plane.” Accessed: Dec. 23, 2023. [Online]. Available: <https://www.materialfoundry.nl/>
- [17] Material Foundry, “Material Plane Github.” Accessed: Dec. 23, 2023. [Online]. Available: <https://github.com/MaterialFoundry/MaterialPlane/wiki/Beta-Hardware-Guide>
- [18] “Robust and Accurate Indoor Localization Using Learning-Based Fusion of Wi-Fi RTT and RSSI.” 2022. Accessed: Dec. 27, 2023. [Online]. Available: <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC9002808/>
- [19] Zhongqin Wang, Ning Ye, Reza Malekian, Fu Xiao, and Ruchuan Wang, “TrackT Accurate tracking of RFID tags with mm-level accuracy using first-order taylor series approximation.” Dec. 15, 2016. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S1570870516302736>
- [20] “YOLOv4 / Scaled-YOLOv4 / YOLO - Neural Networks for Object Detection (Windows and Linux version of Darknet).” Accessed: Dec. 28, 2023. [Online]. Available: <https://github.com/AlexeyAB/darknet>
- [21] Apple, “ARKit.” Accessed: Dec. 27, 2023. [Online]. Available: https://developer.apple.com/documentation/arkit/arkit_in_ios/content_anchors/scanning_and_detecting_3d_objects
- [22] Garrido-Jurado S., Muñoz-Salinas R., Madrid-Cuevas F.J., and Marín-Jiménez M., “Automatic generation and detection of highly reliable fiducial markers under occlusion,” vol. 47, no. 6. Elsevier Science Inc. [Online]. Available: <https://doi.org/10.1016/j.patcog.2014.01.005>
- [23] Zhengyou Zhang and Li-Wei He, “Whiteboard Scanning and Image Enhancement,” vol. 17, no. 2. Elsevier Science Inc., pp. 414–432, Jun. 15, 2006.
- [24] G. Bradski, “Hough Circle Transform.” Accessed: Dec. 27, 2023. [Online]. Available: https://docs.opencv.org/4.x/da/d53/tutorial_py_houghcircles.html
- [25] G. Bradski, “The OpenCV Library.” 2000. Accessed: Dec. 12, 2023. [Online]. Available: <https://github.com/opencv/opencv>
- [26] “DroidCam.” [Online]. Available: <https://www.dev47apps.com/>
- [27] Andrew Zabolotny, “LensFun.” [Online]. Available: <https://lensfun.github.io/lenslist/>
- [28] G. Bradski, “Camera Calibration.” [Online]. Available: https://docs.opencv.org/4.x/dc/dbb/tutorial_py_calibration.html
- [29] “Basic concepts of the homography.” [Online]. Available: https://docs.opencv.org/4.x/d9/dab/tutorial_homography.html#tutorial_homography_Demo2
- [30] Thales Sehn Körting, “Hough Circle Transform.” [Online]. Available: <https://github.com/tkortning/youtube/blob/master/how-circle-hough-transform-works/main.py>
- [31] One Page Rules, “Firefight.” Accessed: Dec. 31, 2023. [Online]. Available: <https://www.onepagerules.com/games/grimdark-future-firefight>